

# MACHINE LEARNING IN HIGH ENERGY PHYSICS

## LECTURE #2



Alex Rogozhnikov, 2015

# RECAPITULATION

- classification, regression
- kNN classifier and regressor
- ROC curve, ROC AUC

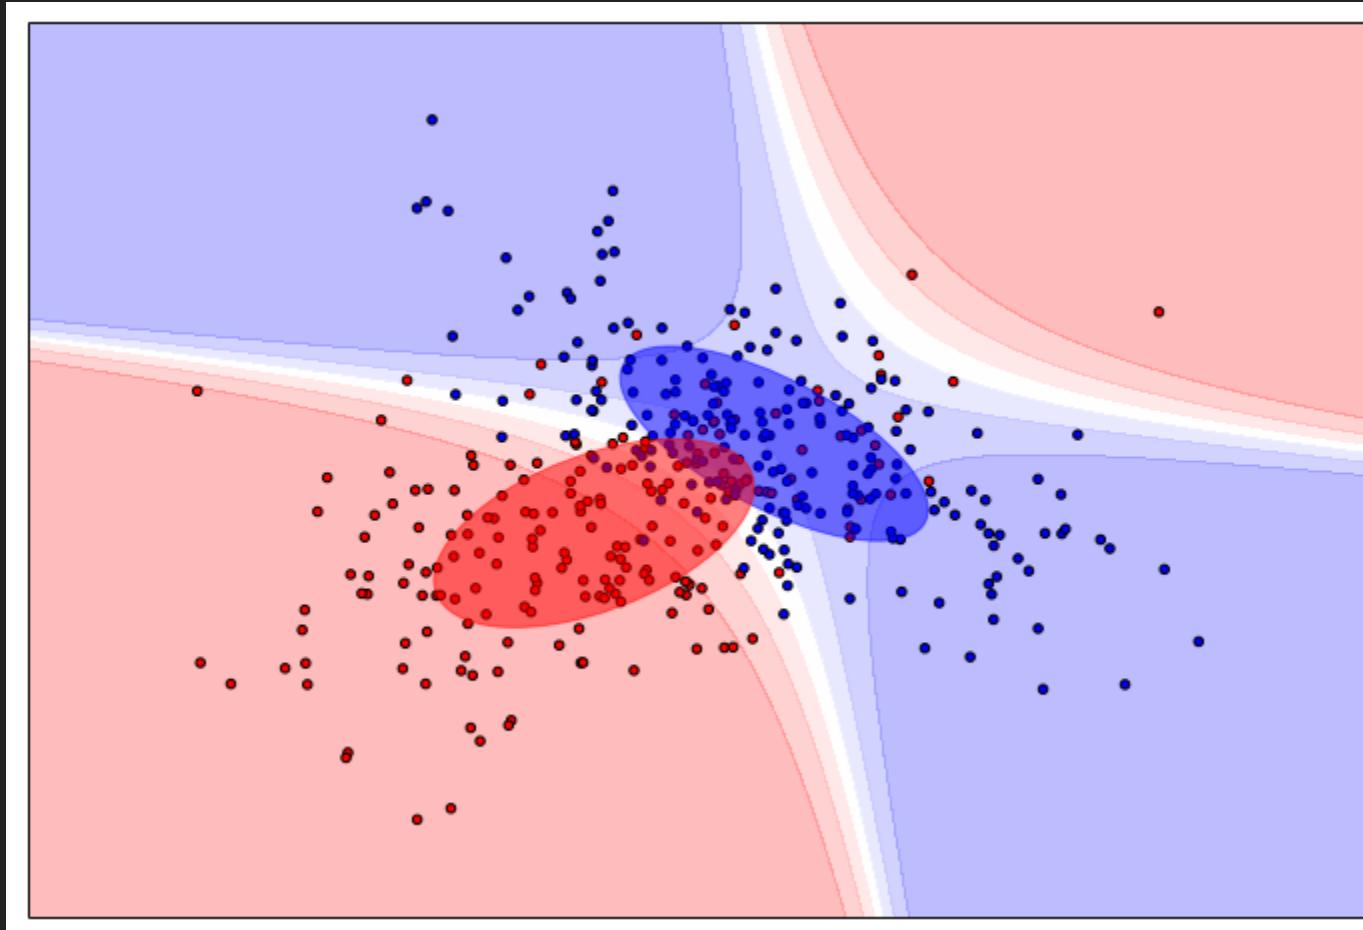
# OPTIMAL BAYESIAN CLASSIFIER

Given knowledge about distributions, we can build optimal classifier

$$\frac{p(y = 1 | x)}{p(y = 0 | x)} = \frac{p(y = 1) p(x | y = 1)}{p(y = 0) p(x | y = 0)}$$

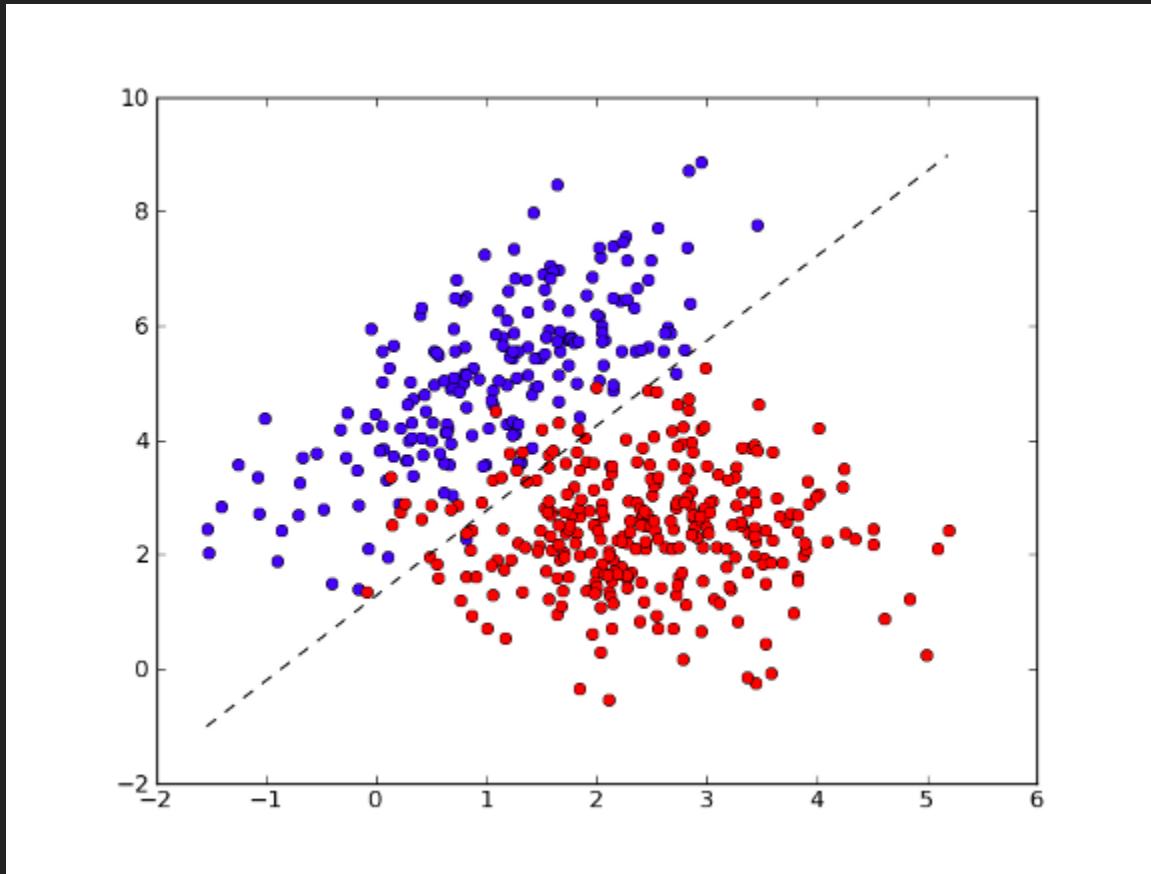
But distributions are **complex**, contain many parameters.

# QDA



QDA follows generative approach.

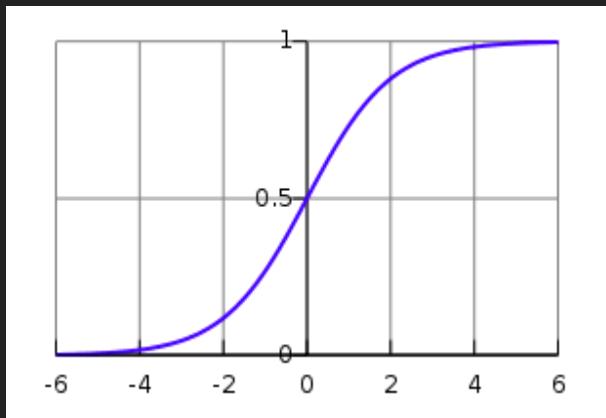
# LOGISTIC REGRESSION



Decision function  $d(x) = \langle w, x \rangle + w_0$

Sharp rule:  $\hat{y} = \text{sgn } d(x)$

# LOGISTIC REGRESSION



$$d(x) = \langle w, x \rangle + w_0$$

Smooth rule:

$$p_{+1}(x) = \sigma(d(x))$$

$$p_{-1}(x) = \sigma(-d(x))$$

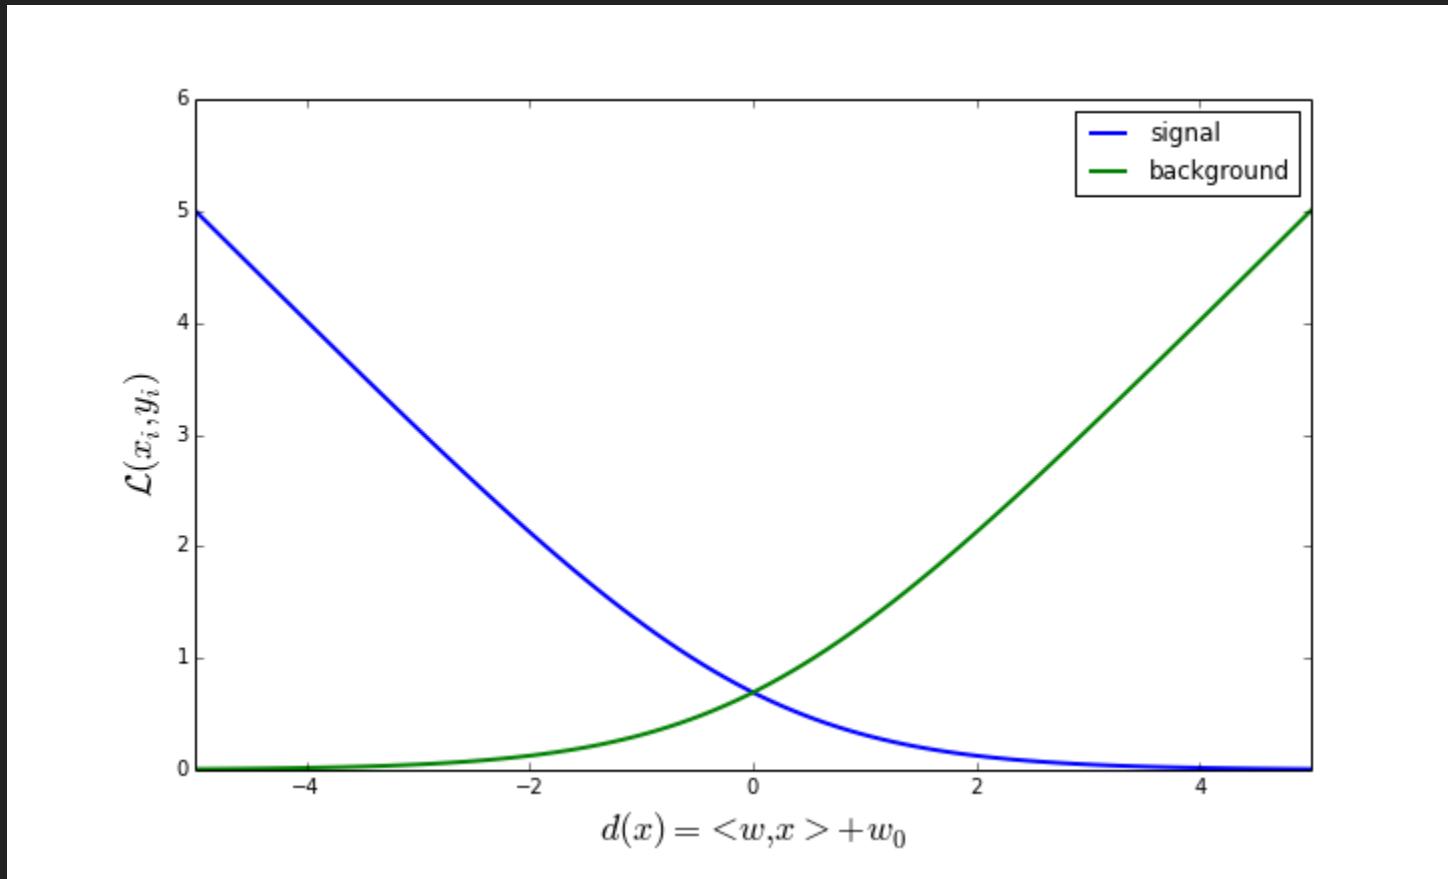
Optimizing weights  $w, w_0$  to maximize log-likelihood

$$\mathcal{L} = \frac{1}{N} \sum_{i \in \text{events}} -\ln(p_{y_i}(x_i)) = \frac{1}{N} \sum_i L(x_i, y_i) \rightarrow \min$$

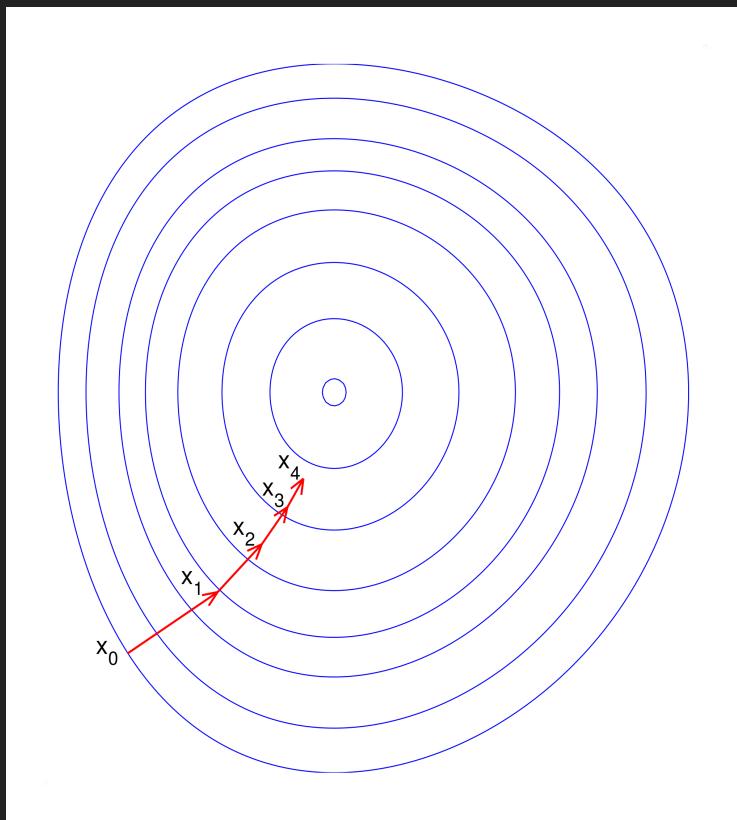
# LOGISTIC LOSS

Loss penalty for single observation

$$L(x_i, y_i) = -\ln(p_{y_i}(x_i)) = \begin{cases} \ln(1 + e^{-d(x_i)}), & y_i = +1 \\ \ln(1 + e^{d(x_i)}), & y_i = -1 \end{cases}$$



# GRADIENT DESCENT & STOCHASTIC OPTIMIZATION



Problem:  
finding  $w$  to minimize  $\mathcal{L}$

$$w \leftarrow w - \eta \frac{\partial \mathcal{L}}{\partial w}$$

$\eta$  is step size  
(also `shrinkage`, `learning rate`)

# STOCHASTIC GRADIENT DESCENT

$$\mathcal{L} = \frac{1}{N} \sum_i L(x_i, y_i) \rightarrow \min$$

On each iteration make a step with respect to only one event:

1. take  $i$  – random event from training data
2.  $w \leftarrow w - \eta \frac{\partial \mathcal{L}(x_i, y_i)}{\partial w}$

Each iteration is done much faster, but training process is less stable.

# POLYNOMIAL DECISION RULE

$$d(x) = w_0 + \sum_i w_i x_i + \sum_{ij} w_{ij} x_i x_j$$

is again linear model, introduce new features:

$$z = \{1\} \cup \{x_i\}_i \cup \{x_i x_j\}_{ij}$$

$$d(x) = \sum_i w_i z_i$$

and reusing logistic regression.

We can add  $x_0 = 1$  as one more variable to dataset and forget about intercept

$$d(x) = w_0 + \sum_{i=1}^N w_i x_i = \sum_{i=0}^N w_i x_i$$

# PROJECTING IN HIGHER DIMENSION SPACE

SVM with polynomial kernel visualization



After adding new features, classes may become separable.

# KERNEL TRICK

$P$  is projection operator (which adds new features).

$$d(x) = \langle w, P(x) \rangle$$

Assume

$$w = \sum_i \alpha_i P(x_i)$$

and look for optimal  $\alpha_i$

$$d(x) = \sum_i \alpha_i \langle P(x_i), P(x) \rangle = \sum_i \alpha_i K(x_i, x)$$

We need only kernel:  $K(x, y) = \langle P(x), P(y) \rangle$

# KERNEL TRICK

Popular kernel is gaussian Radial Basis Function:

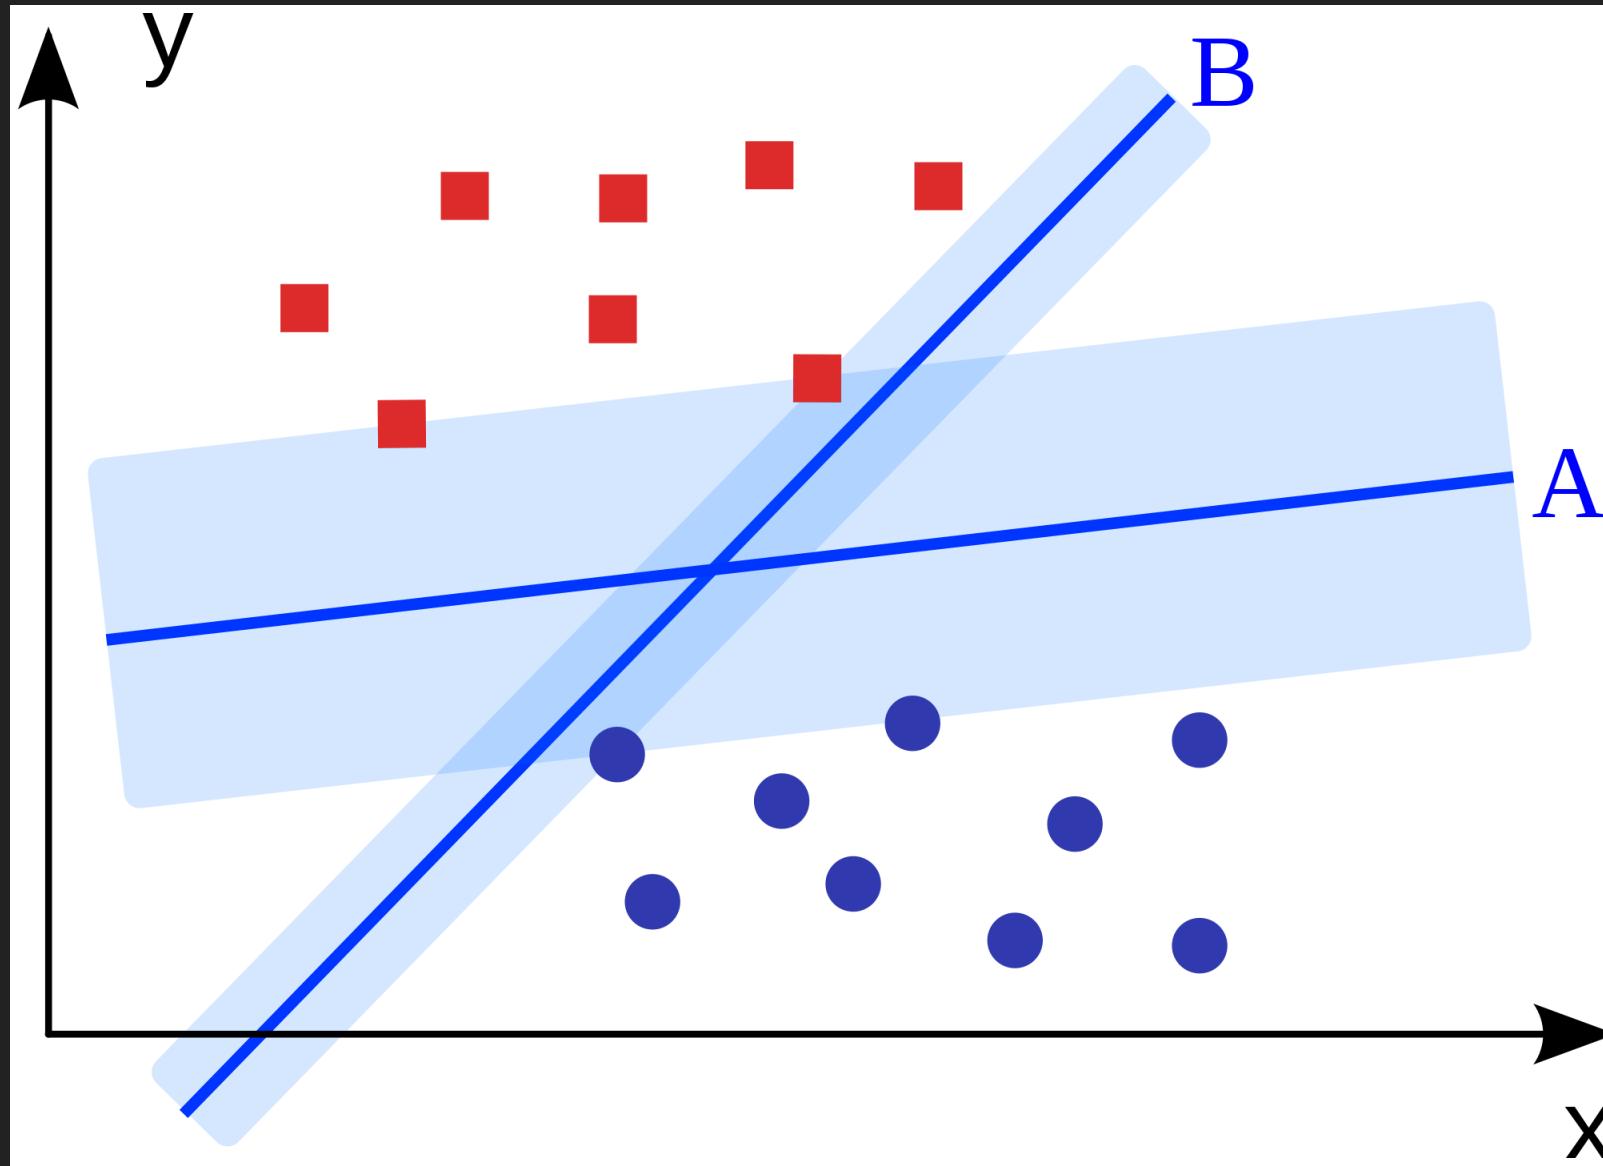
$$K(x, y) = e^{-c\|x-y\|^2}$$

Corresponds to projection to Hilbert space.

Exercise: find a correspond projection.

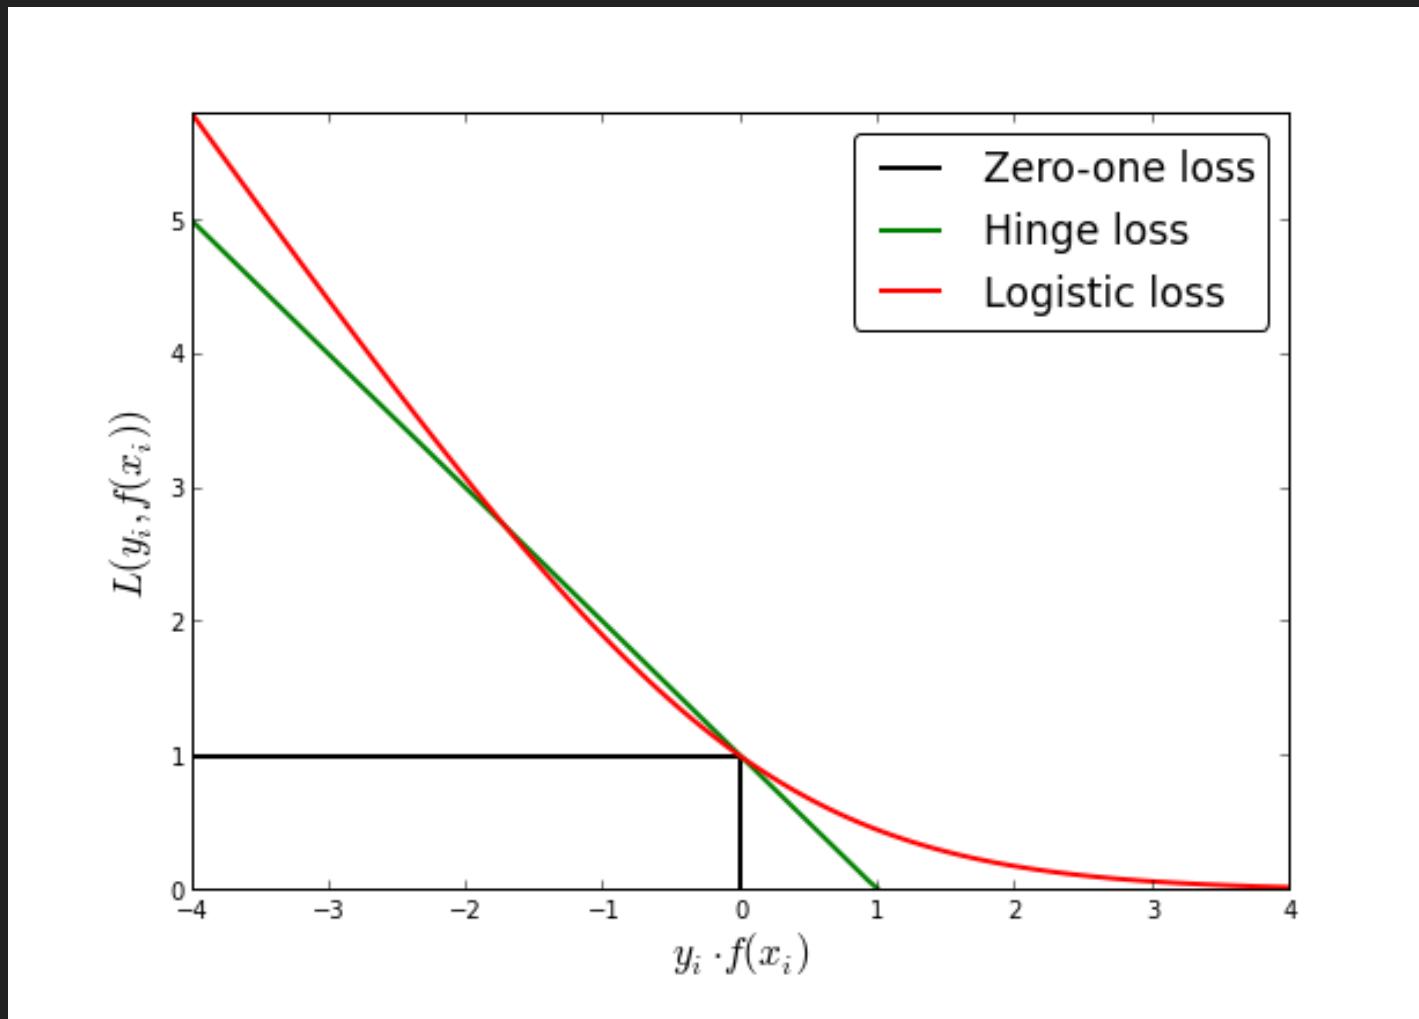
# SUPPORT VECTOR MACHINE

SVM selects decision rule with maximal possible margin.

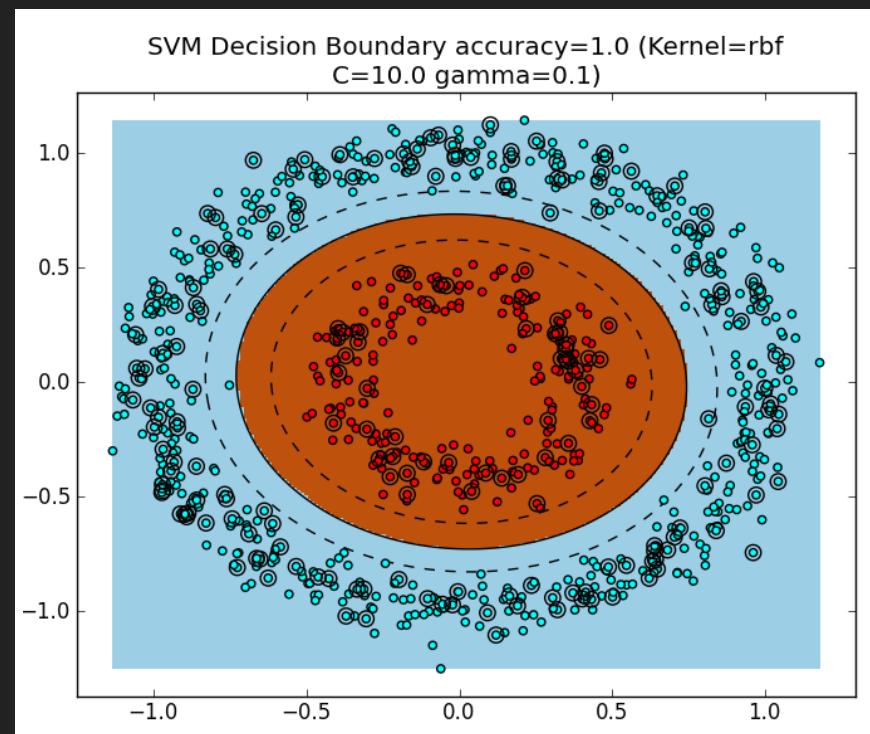
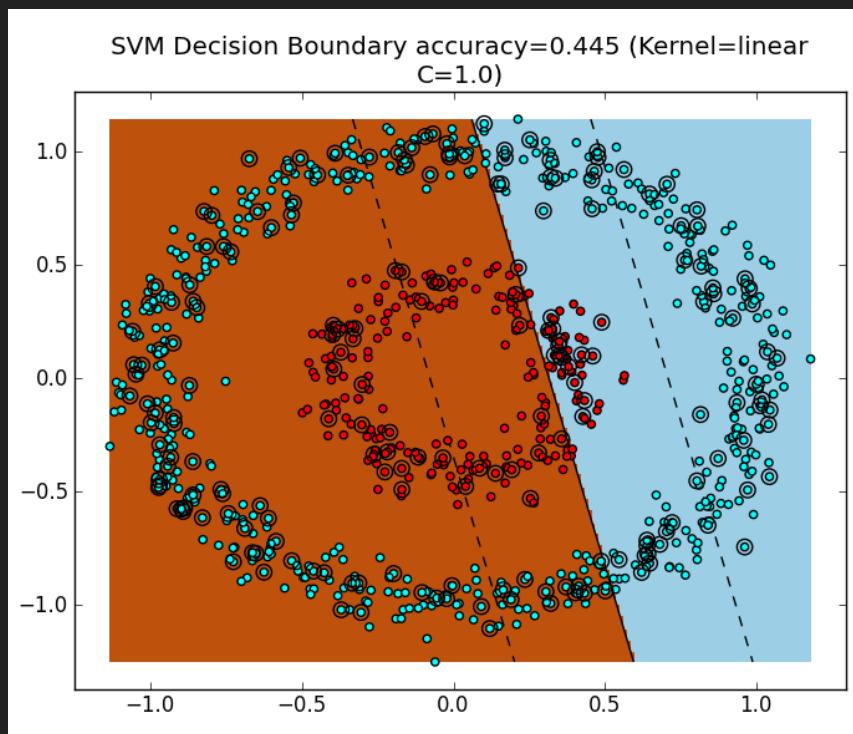


# HINGE LOSS FUNCTION

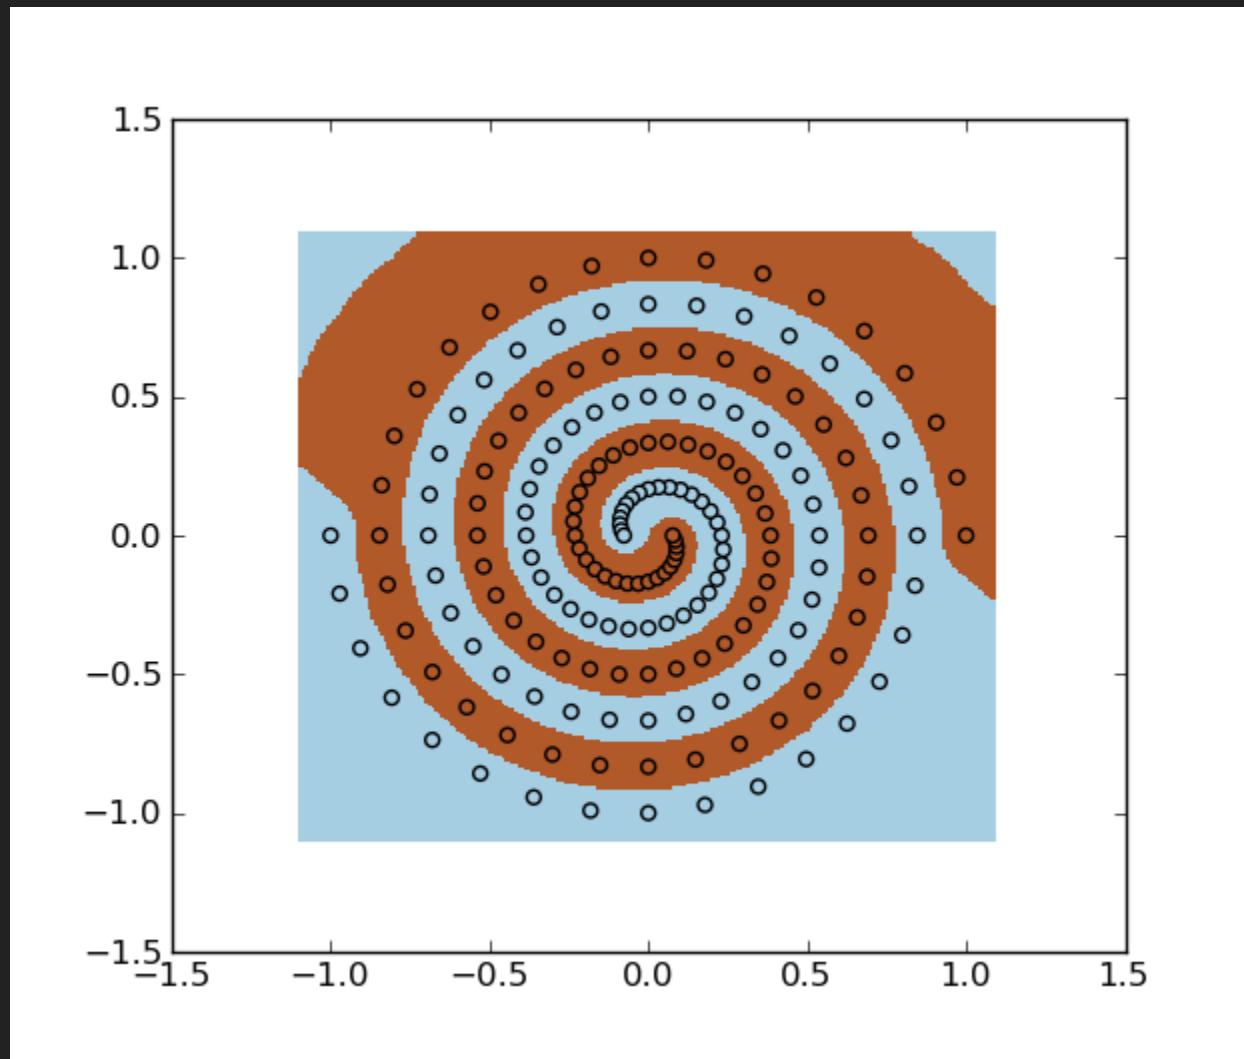
SVM uses different loss function (only signal losses compared):



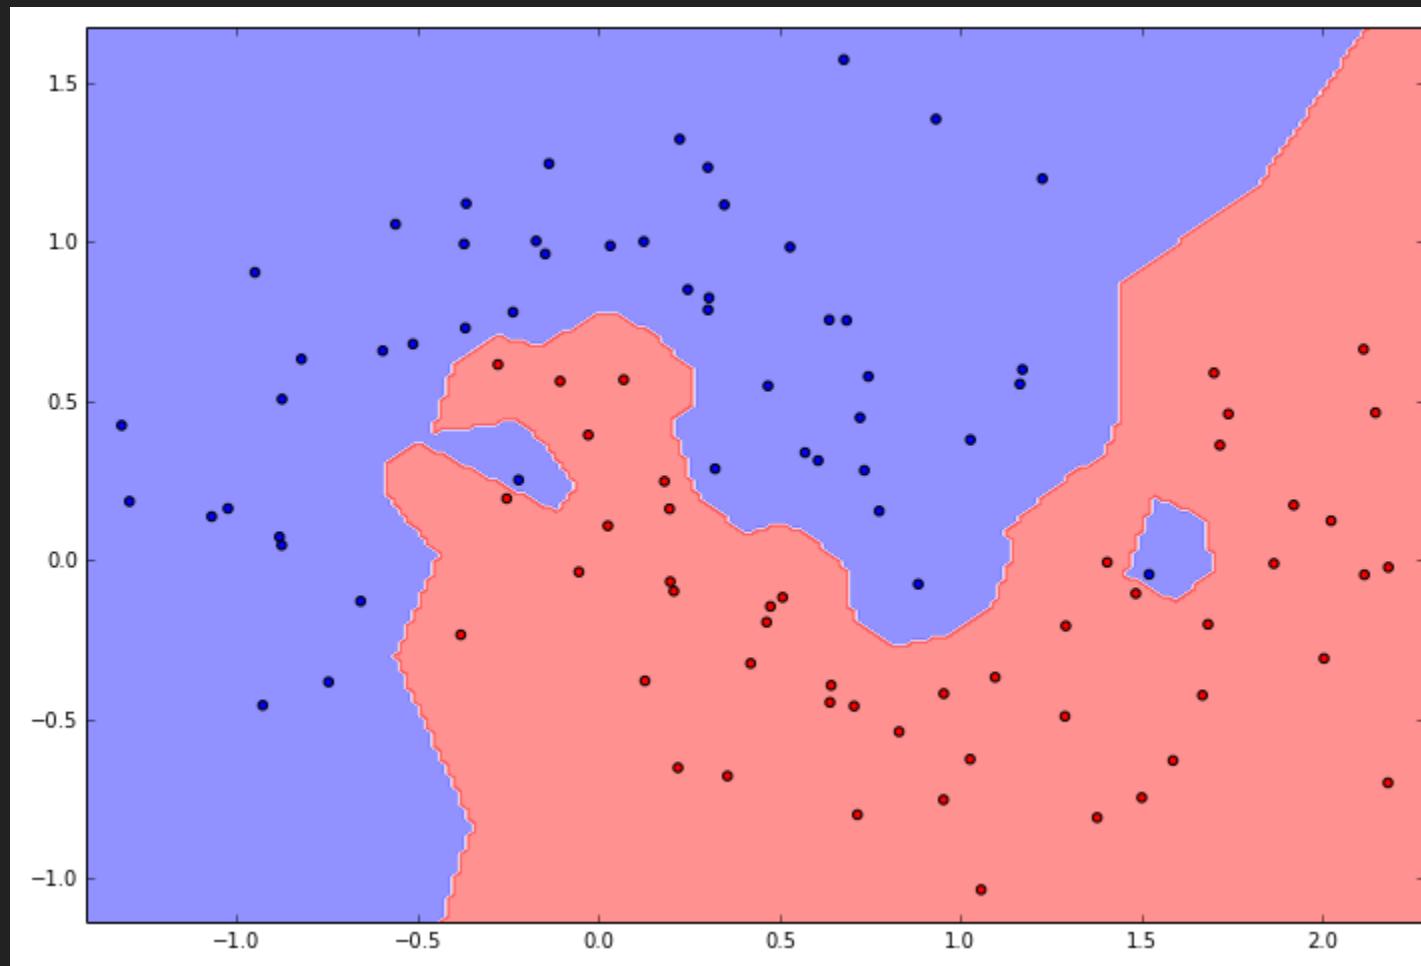
# SVM + RBF KERNEL



# SVM + RBF KERNEL



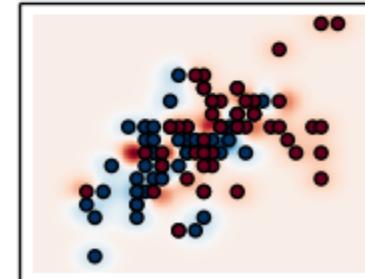
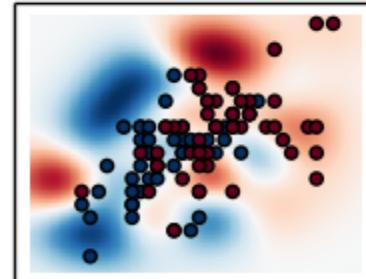
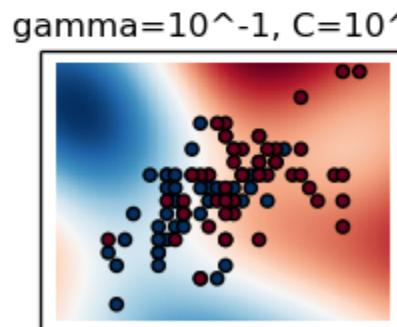
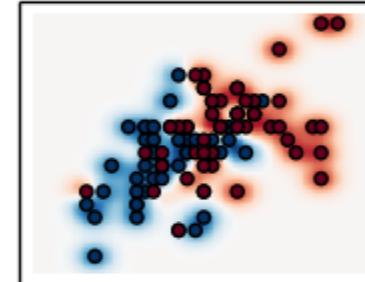
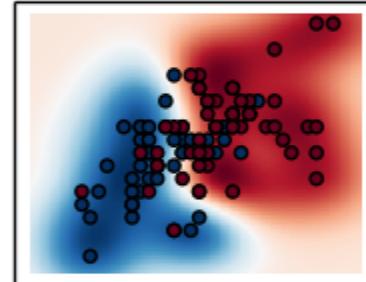
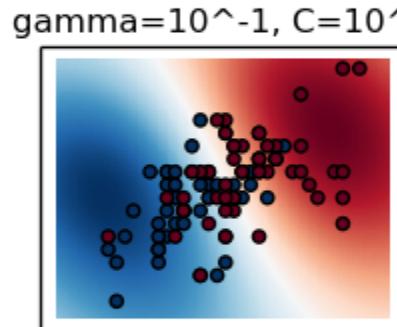
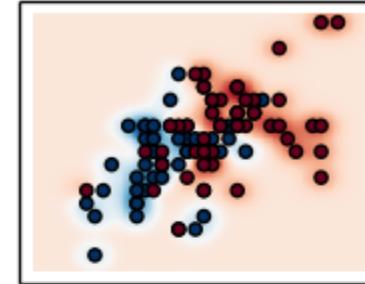
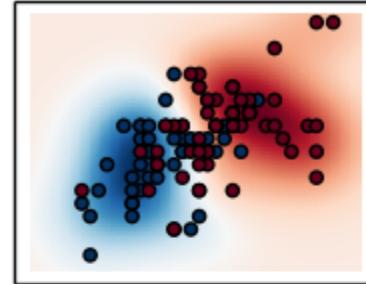
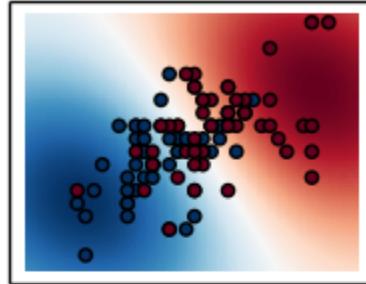
# OVERFITTING



Knn with  $k=1$  gives ideal classification of training data.

# OVERFITTING

gamma=10<sup>-1</sup>, C=10<sup>-2</sup> gamma=10<sup>0</sup>, C=10<sup>-2</sup> gamma=10<sup>1</sup>, C=10<sup>-2</sup>



There are two definitions of overfitting, which often coincide.

## DIFFERENCE-OVERFITTING

There is significant difference in quality of predictions between train and test.

## COMPLEXITY-OVERFITTING

Formula has too high complexity (e.g. too many parameters), increasing the number of parameters drives to lower quality.

# MEASURING QUALITY

To get unbiased estimate, one should test formula on independent samples (and be sure that no train information was given to algorithm during training)

In most cases, simply splitting data into train and holdout is enough.

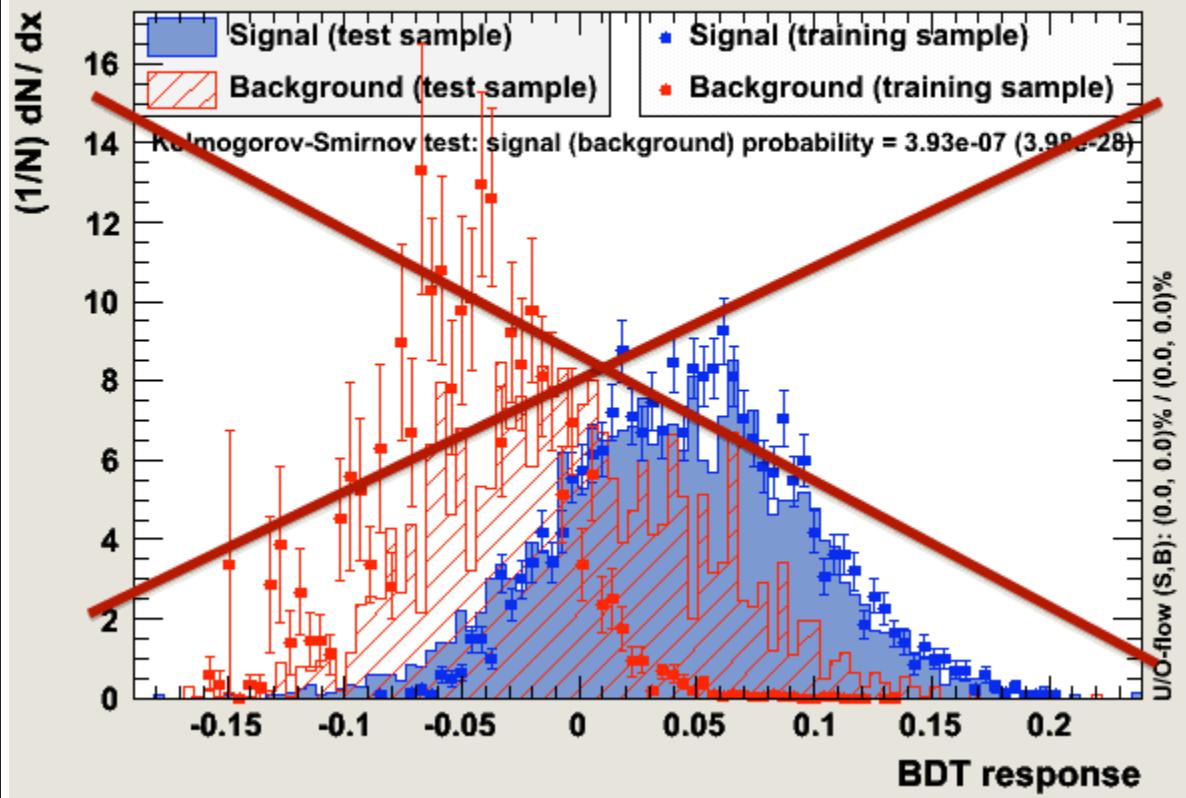
More approaches in seminar.

**Difference-overfitting** is inessential, provided that we measure quality on holdout (though easy to check).

**Complexity-overfitting** is problem — we need to test different parameters for optimality (more examples through the course).

Don't use distribution comparison to detect overfitting

# TMVA overtraining check for classifier: BDT



# REGULARIZATION

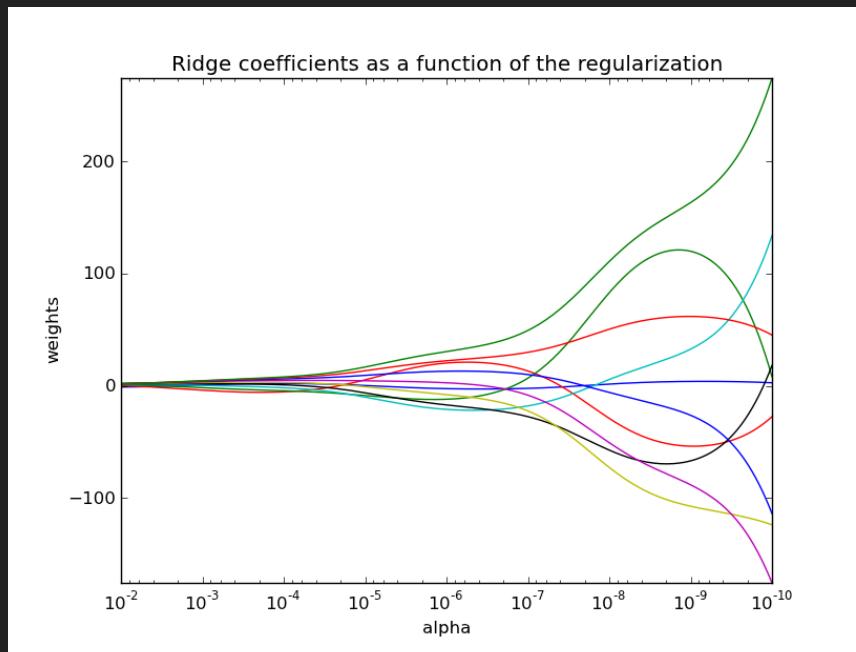
When number of weights is high, overfitting is very probable

Adding regularization term to loss function:

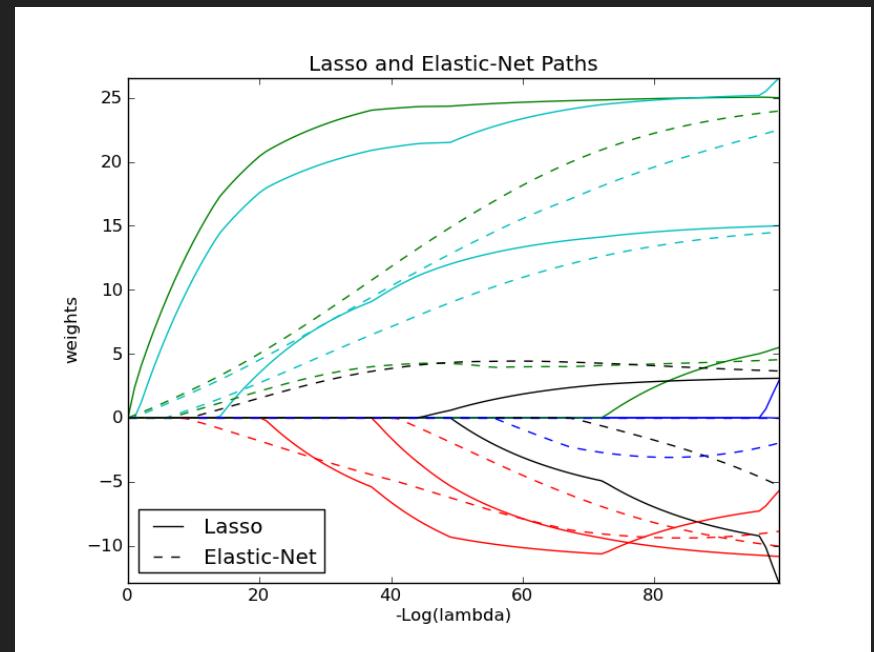
$$\mathcal{L} = \frac{1}{N} \sum_i L(x_i, y_i) + \mathcal{L}_{\text{reg}} \rightarrow \min$$

- $L_2$  regularization :  $\mathcal{L}_{\text{reg}} = \alpha \sum_j |w_j|^2$
- $L_1$  regularization:  $\mathcal{L}_{\text{reg}} = \beta \sum_j |w_j|$
- $L_1 + L_2$  regularization:  $\mathcal{L}_{\text{reg}} = \alpha \sum_j |w_j|^2 + \beta \sum_j |w_j|$

# $L_2$ , $L_1$ – REGULARIZATIONS



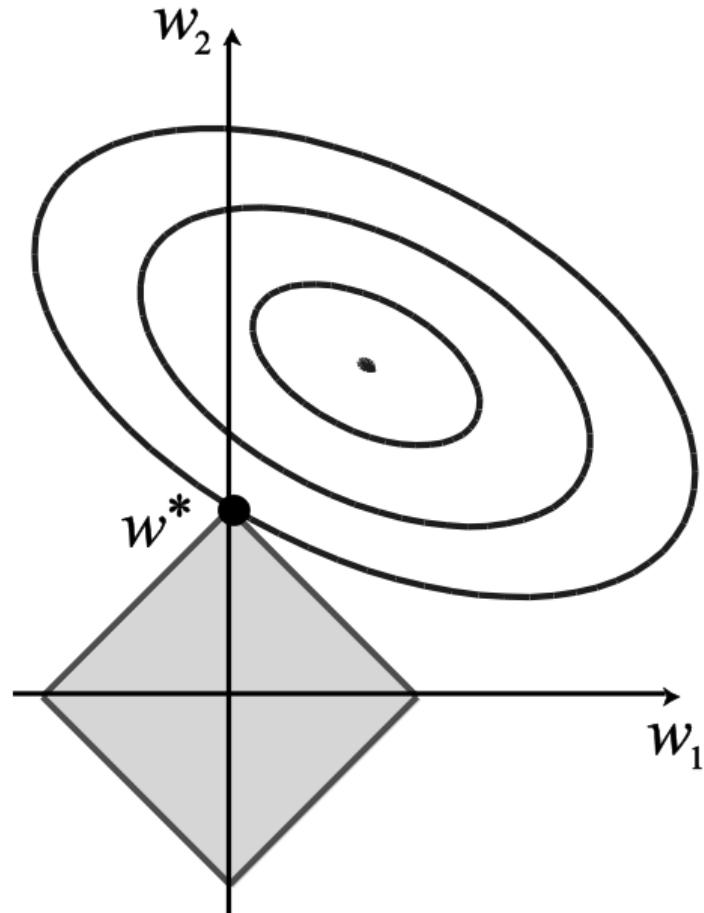
$L_2$  regularization



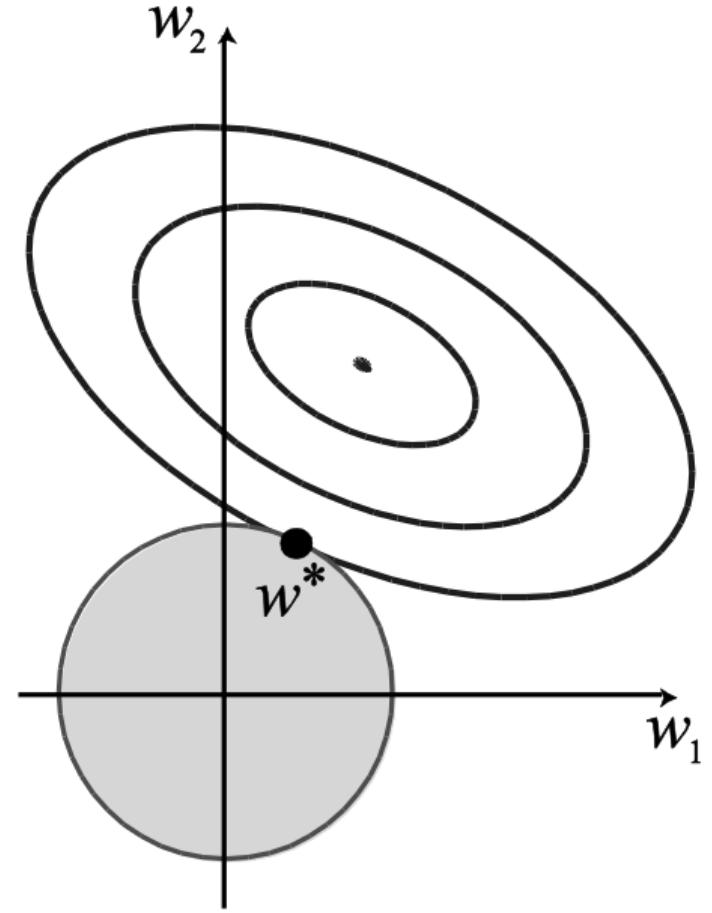
$L_1$  (solid),  $L_1 + L_2$  (dashed)

# REGULARIZATIONS

$L_1$  regularization encourages sparsity

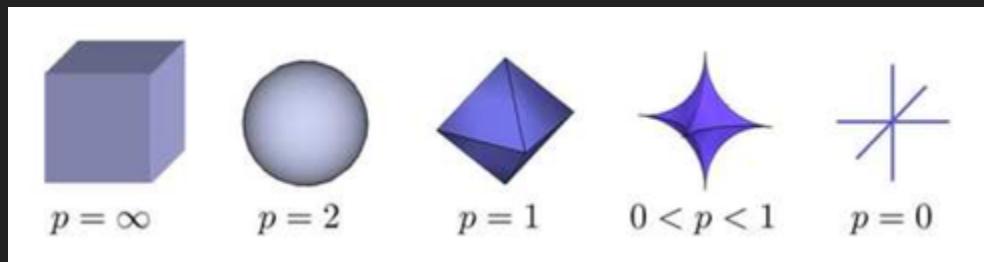


L1



L2

# $L_p$ REGULARIZATIONS



$$L_p = \sum_i w_i^p$$

- What is the expression for  $L_0$ ?
- $L_0 = \sum_i [w_i \neq 0]$   
But nobody uses it, even  $L_p$ ,  $0 < p < 1$ . Why?
- Because it is not convex

# LOGISTIC REGRESSION

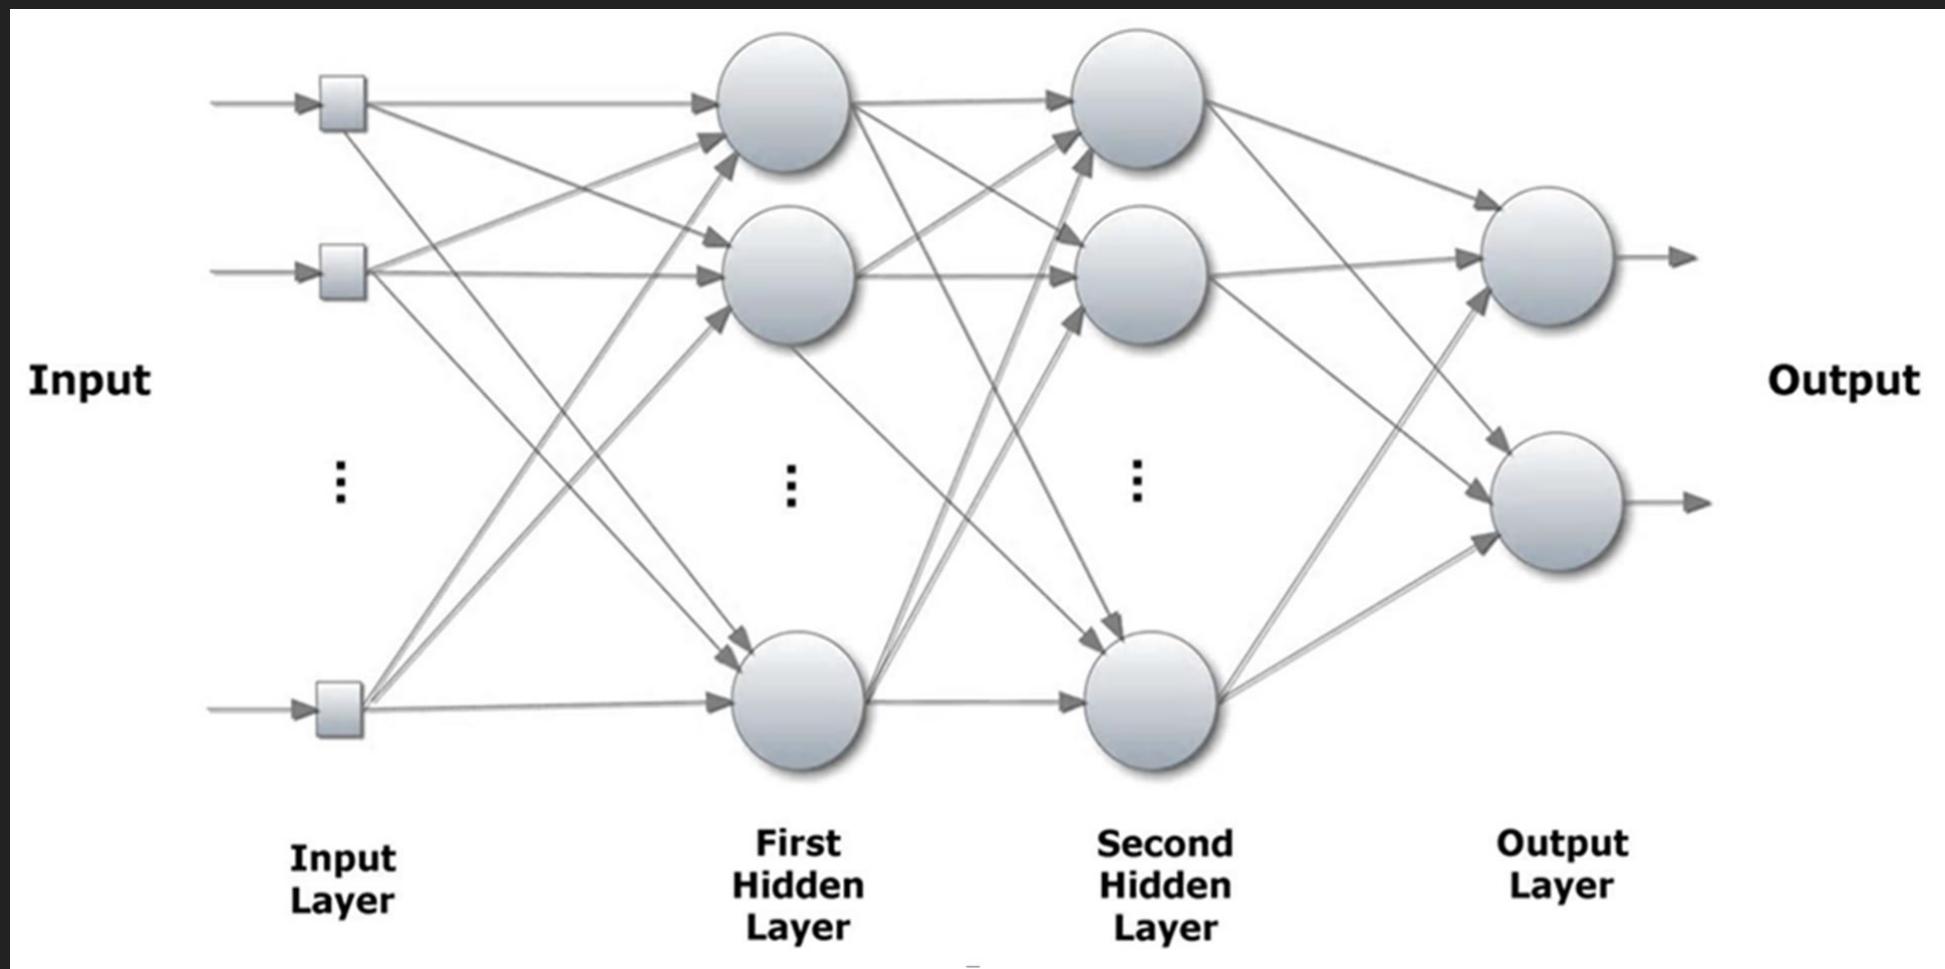
- classifier based on linear decision rule
- training is reduced to convex optimization
- other decision rules are achieved by adding new features
- stochastic optimization is used
- can handle  $> 1000$  features, requires regularization
- no interaction between features

# [ARTIFICIAL] NEURAL NETWORKS

Based on our understanding of natural neural networks

- neurons are organized in networks
- receptors activate some neurons, neurons are activating other neurons, etc.
- connection is via synapses

# STRUCTURE OF ARTIFICIAL FEED-FORWARD NETWORK



# ACTIVATION OF NEURON

Neuron states:  $n = \begin{cases} 1, & \text{activated} \\ 0, & \text{not activated} \end{cases}$

Let  $n_i$  to be state of  $w_i$  to be weight of connection between  $i$ -th neuron and output neuron:

$$n = \begin{cases} 1, & \sum_i w_i n_i > 0 \\ 0, & \text{otherwise} \end{cases}$$

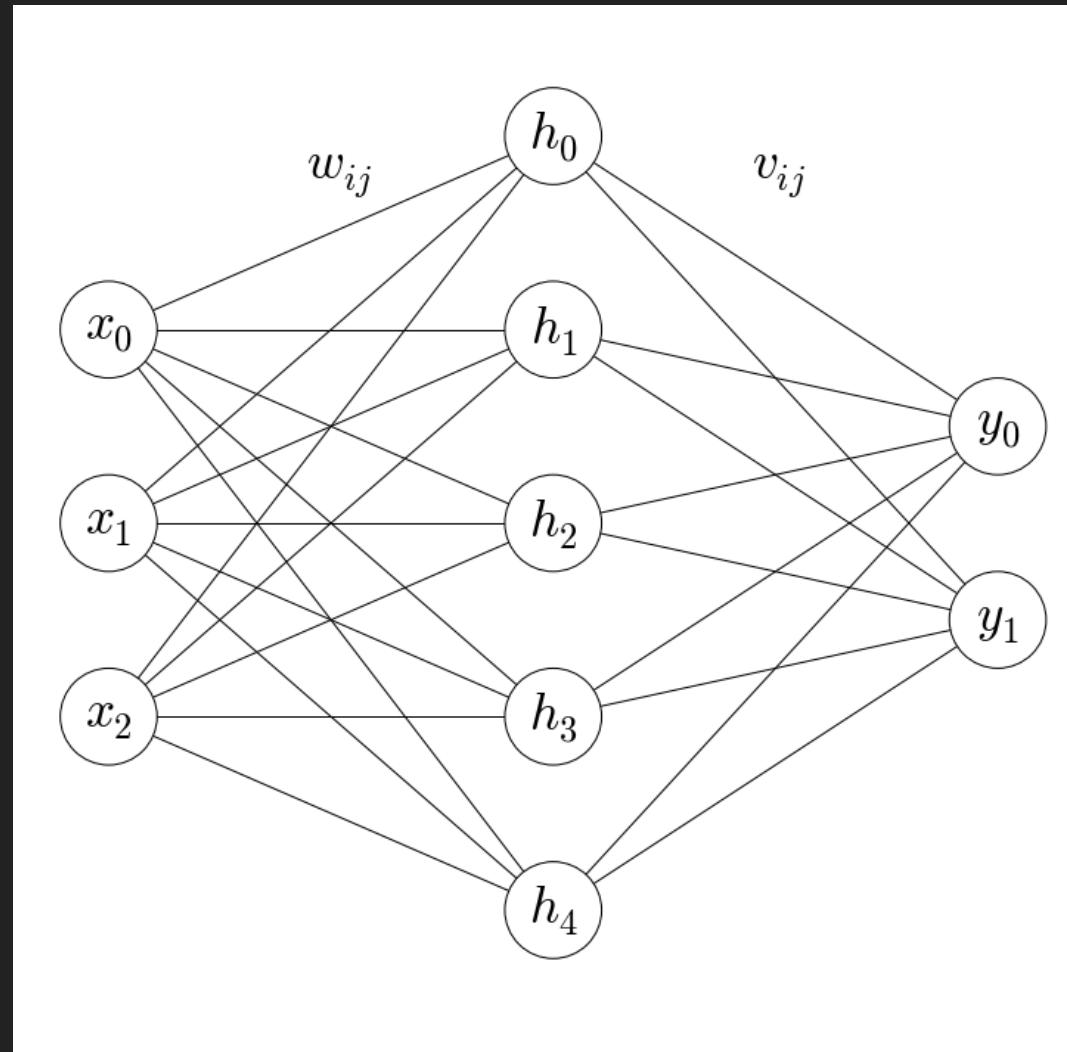
Problem: find set of weights, that minimizes error on train dataset. (discrete optimization)

# SMOOTH ACTIVATIONS:

ONE HIDDEN  
LAYER

$$h_i = \sigma\left(\sum_j w_{ij}x_j\right)$$

$$y_i = \sigma\left(\sum_i v_{ij}h_j\right)$$



# VISUALIZATION OF NN

# NEURAL NETWORKS

- Powerful general purpose algorithm for classification and regression
- Non-interpretable formula
- Optimization problem is **non-convex** with local optimums and has many parameters  
Stochastic optimization speeds up process and helps not to be caught in local minimum.
- **Overfitting** due to large amount of parameters  
 $L_1, L_2$  – regularizations (and other tricks)

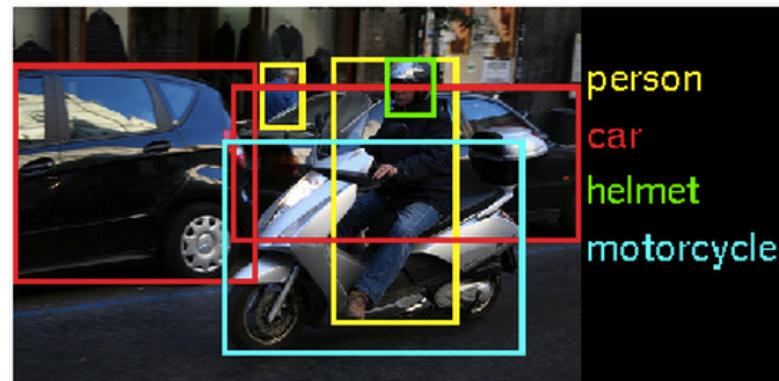
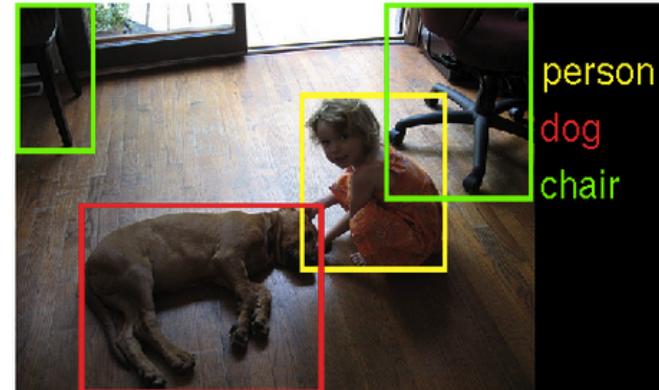
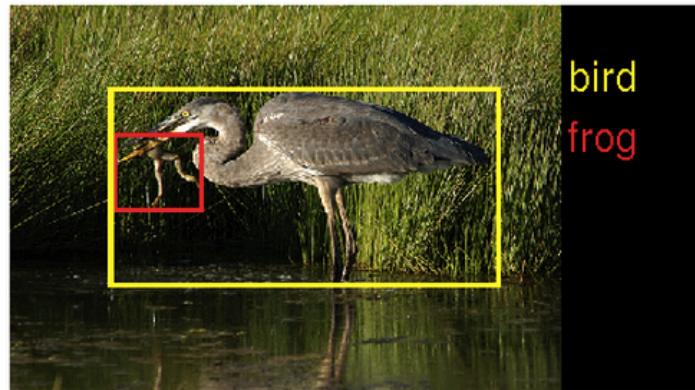
X MINUTES BREAK

# DEEP LEARNING

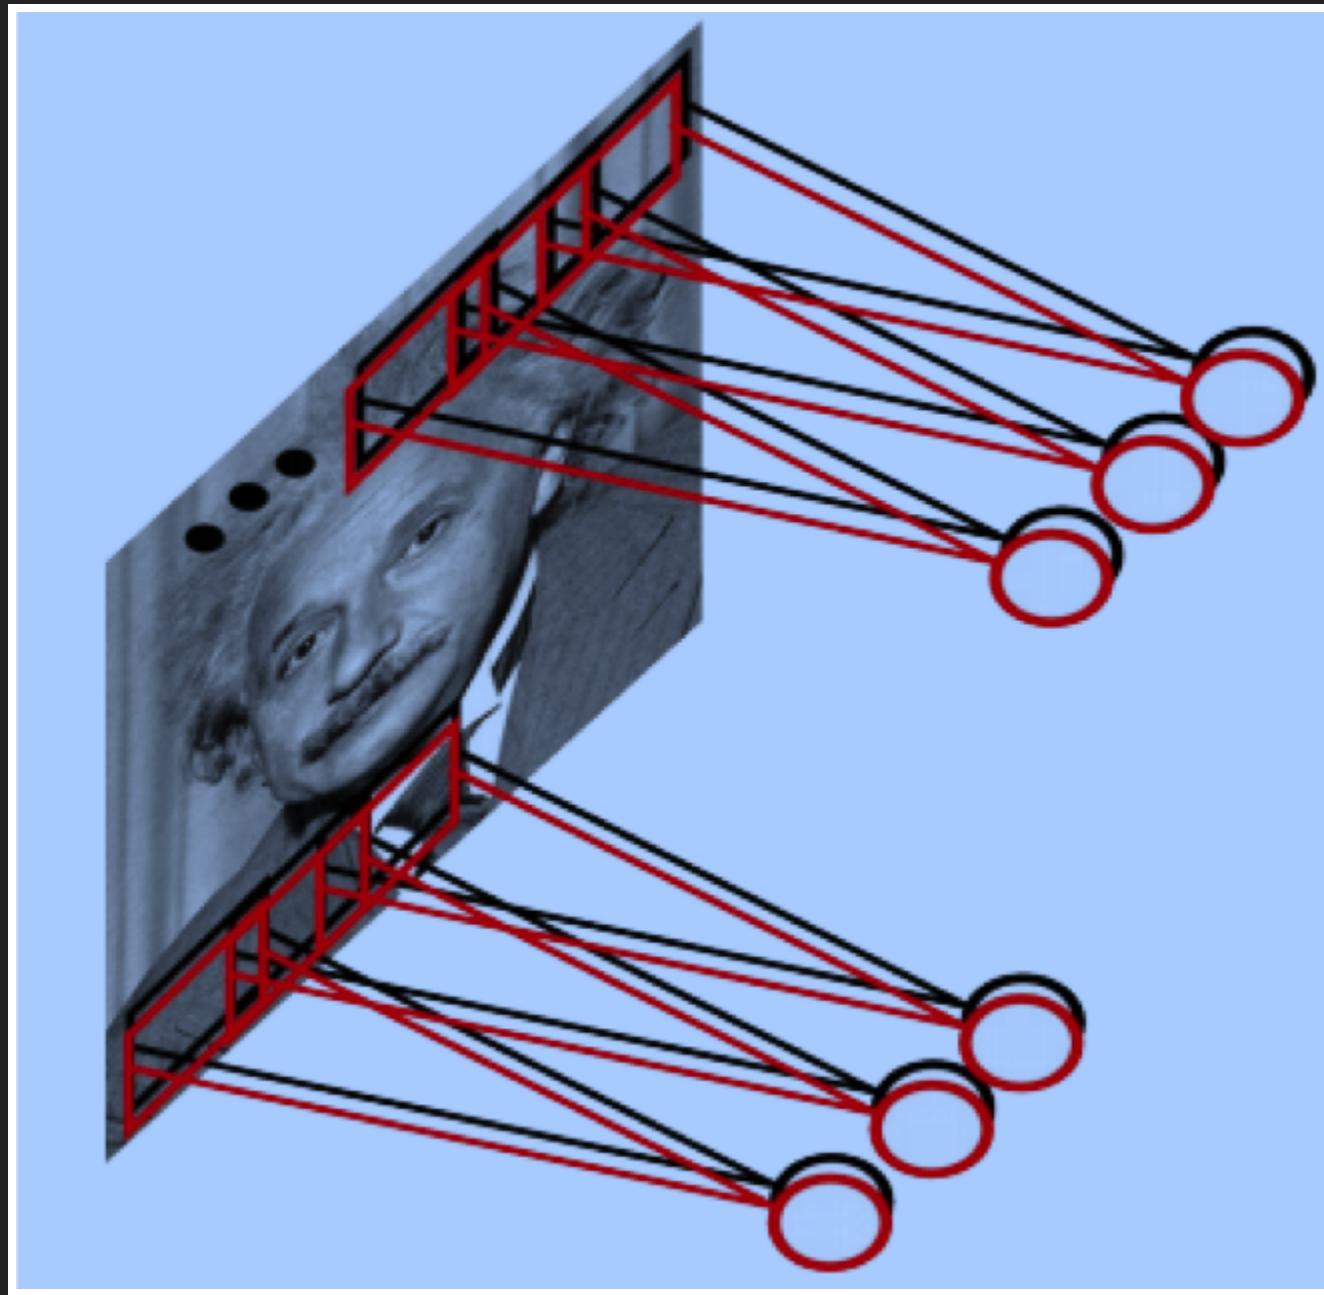
Gradient **diminishes** as number of hidden layers grows.  
Usually 1-2 hidden layers are used.

But modern ANN for image recognition have 7-15 layers.

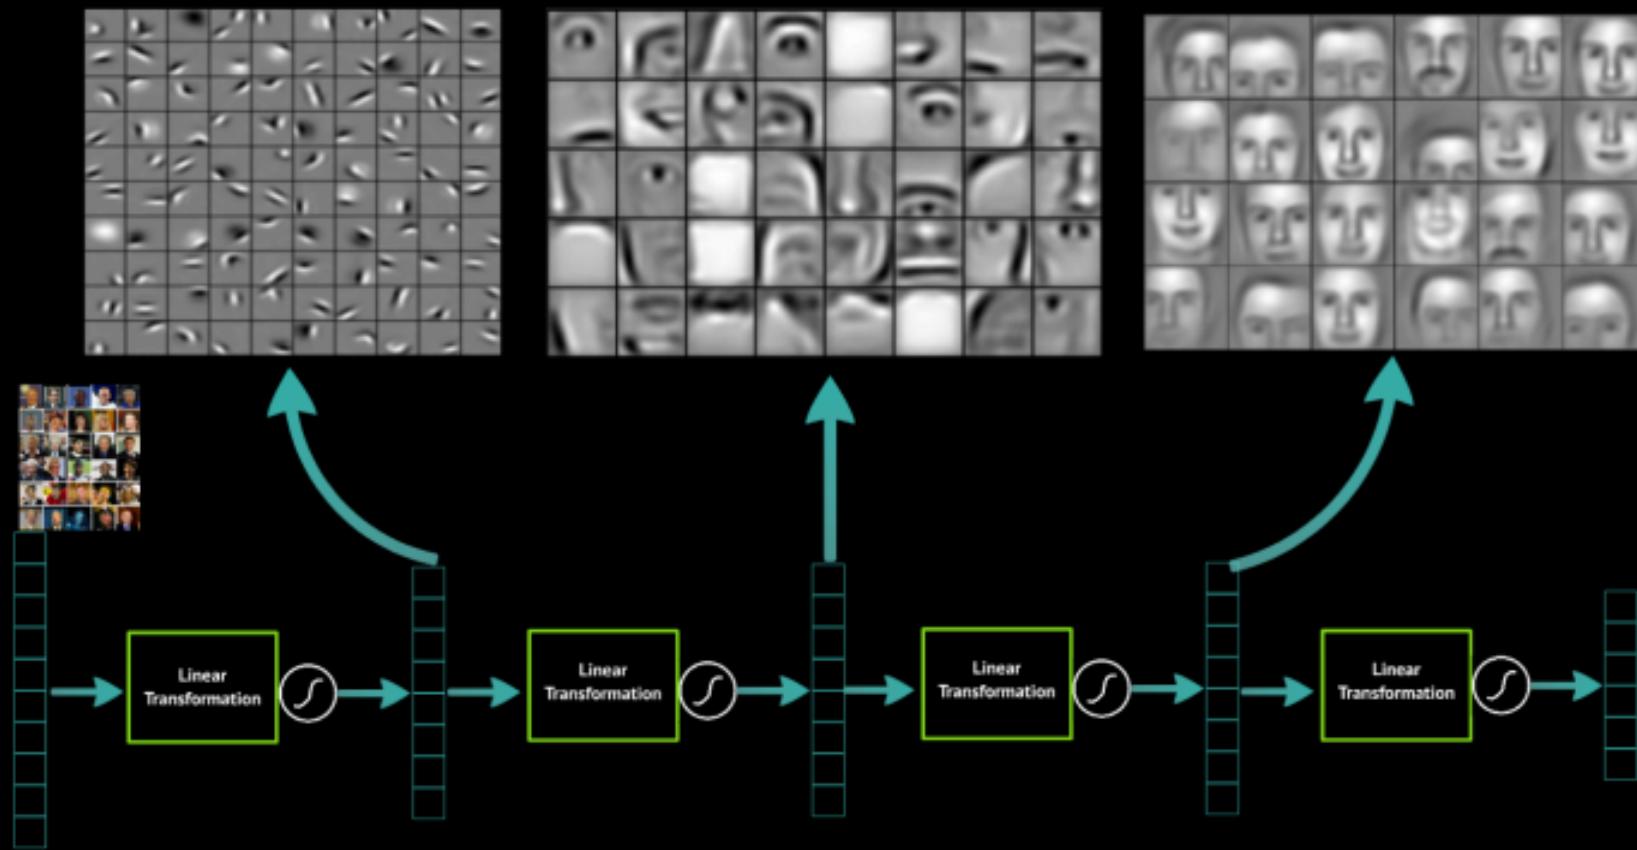
Example ILSVRC2014 images:



# CONVOLUTIONAL NEURAL NETWORK

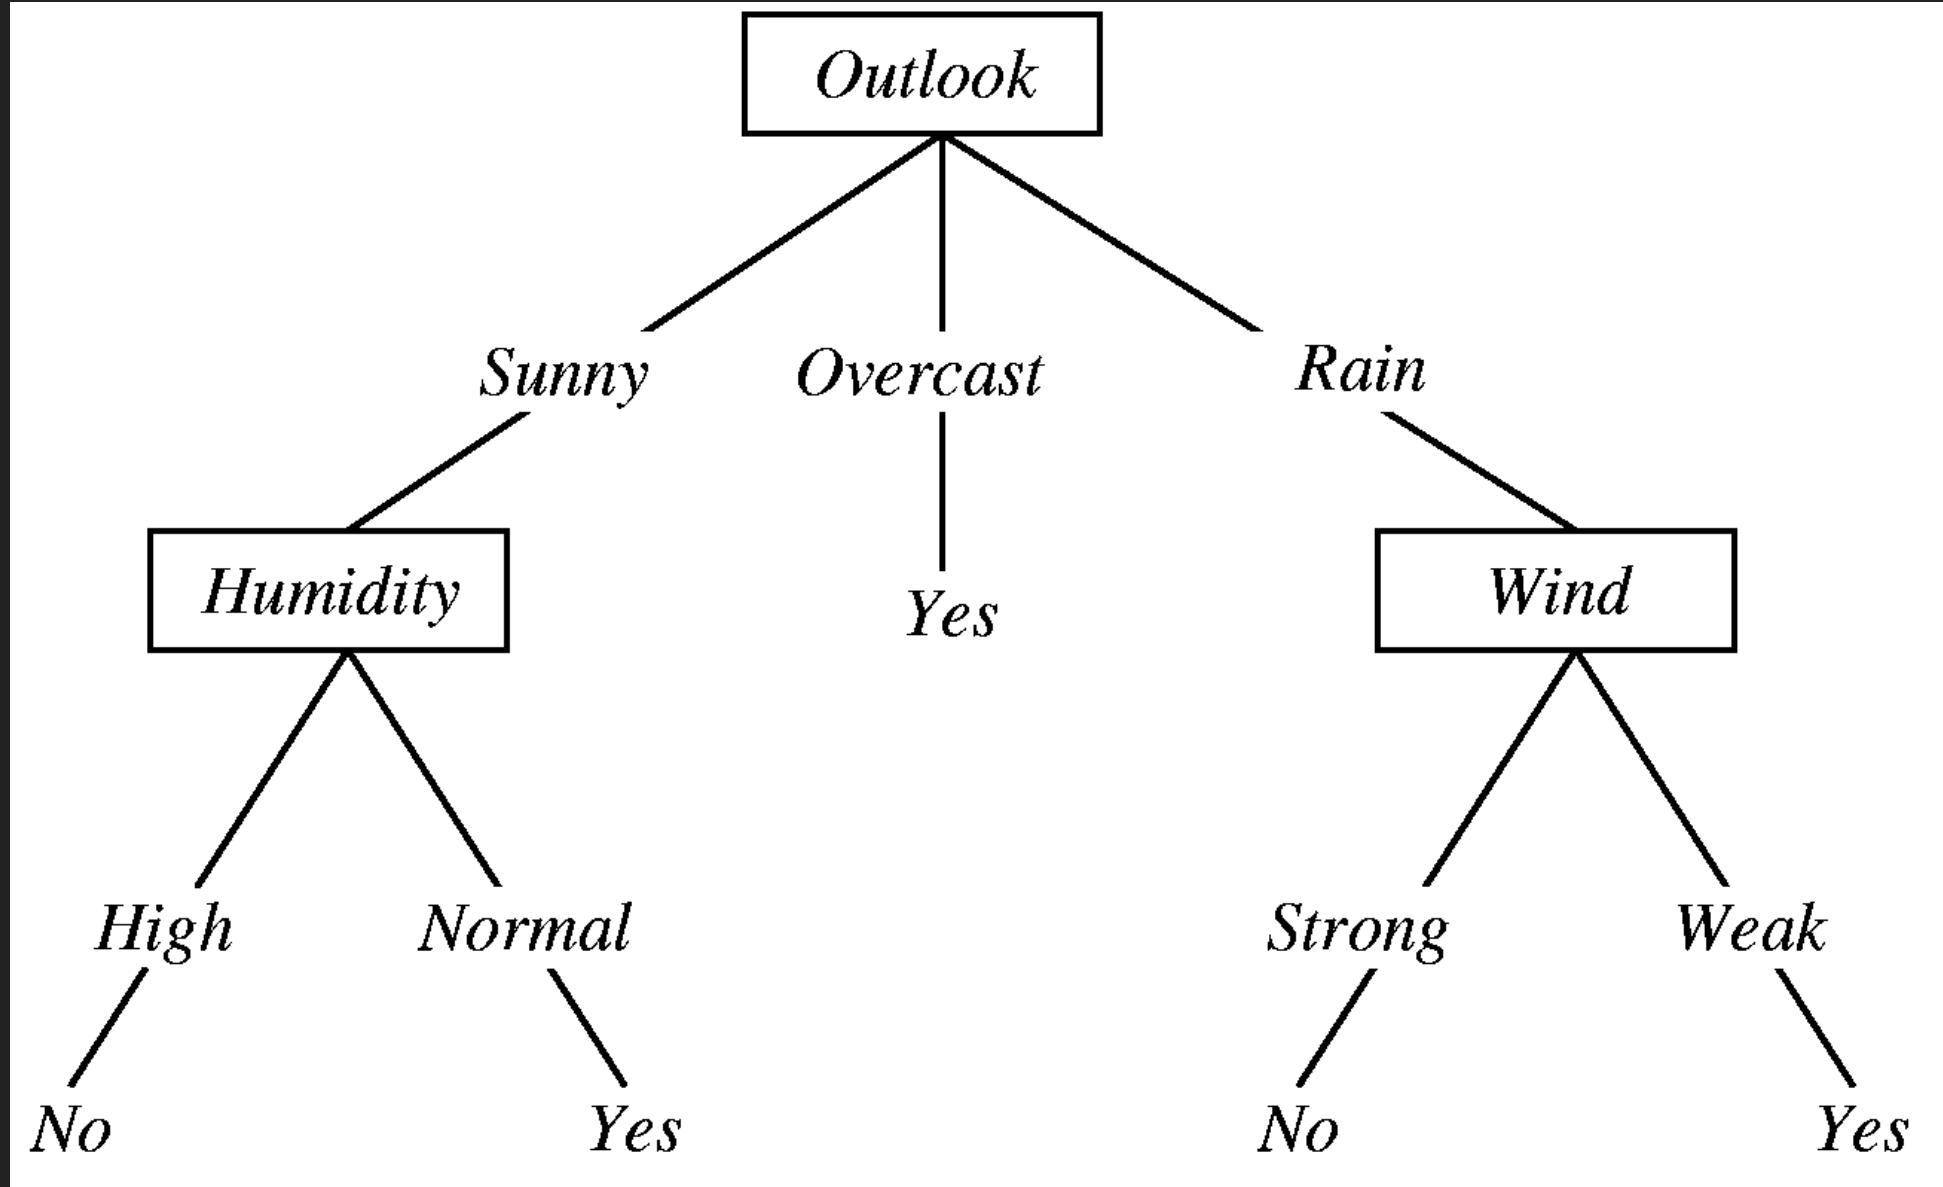


# Deep Learning learns layers of features

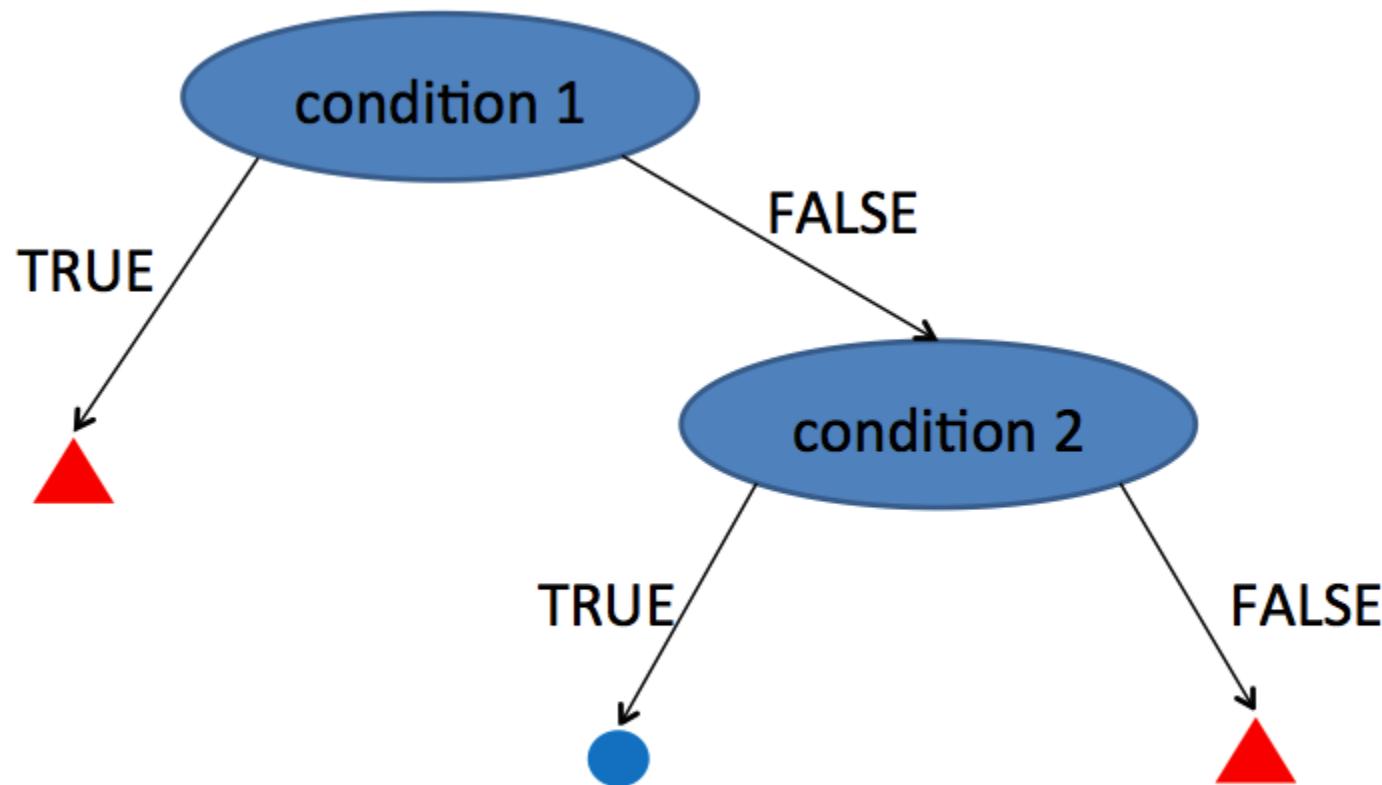


# DECISION TREES

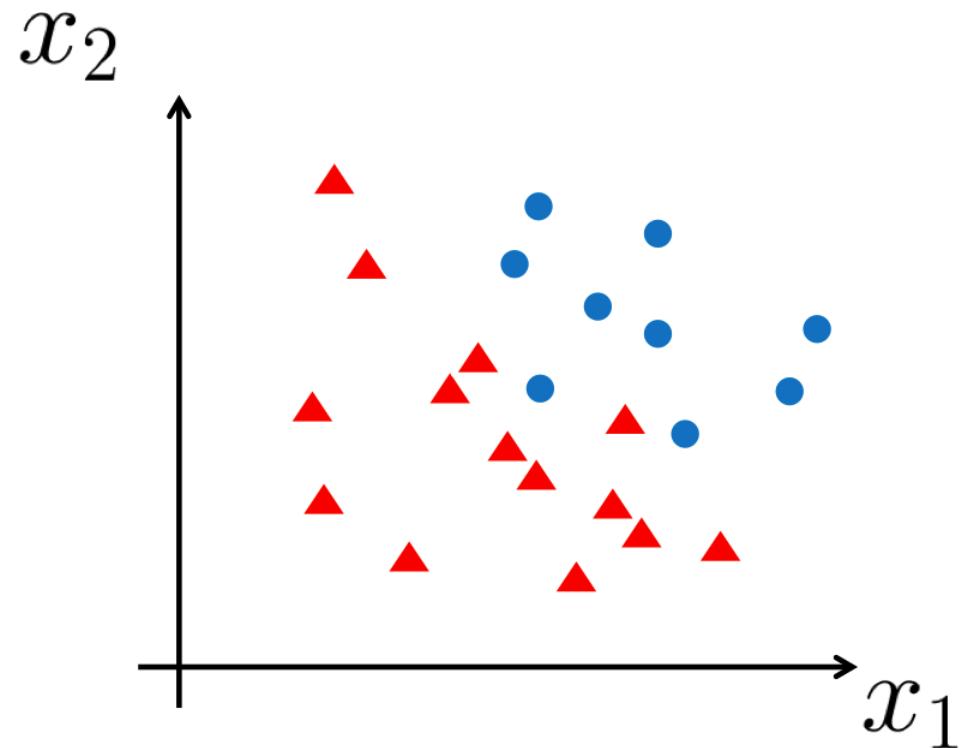
Example: predict outside play based on weather conditions.



# DECISION TREES: IDEA

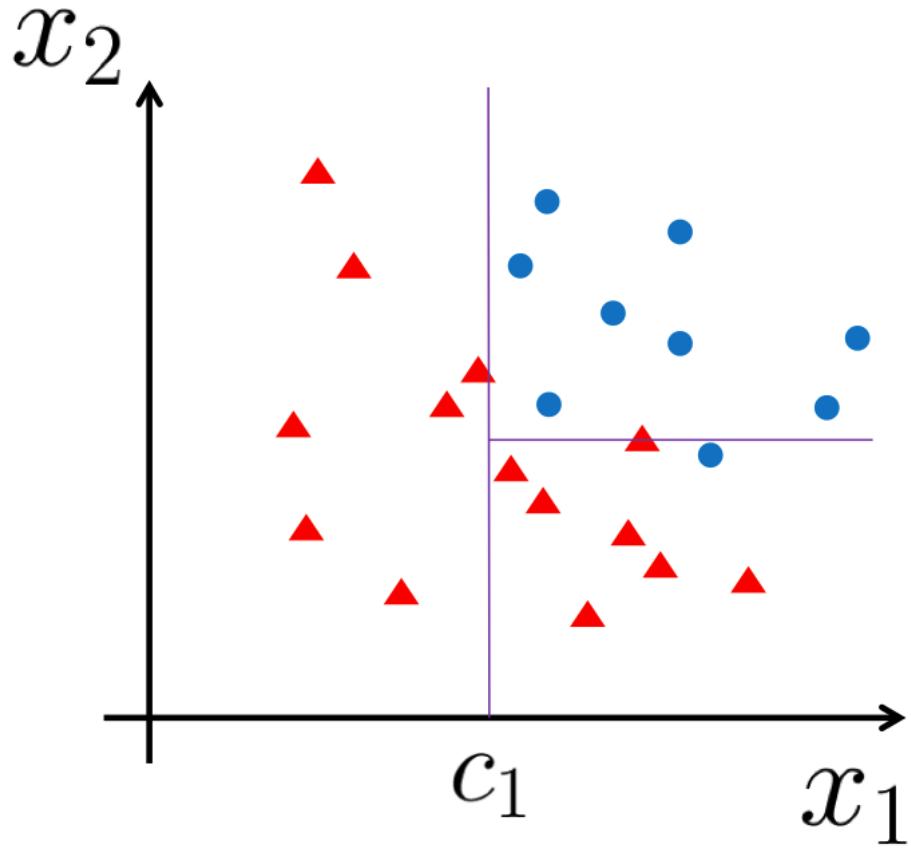
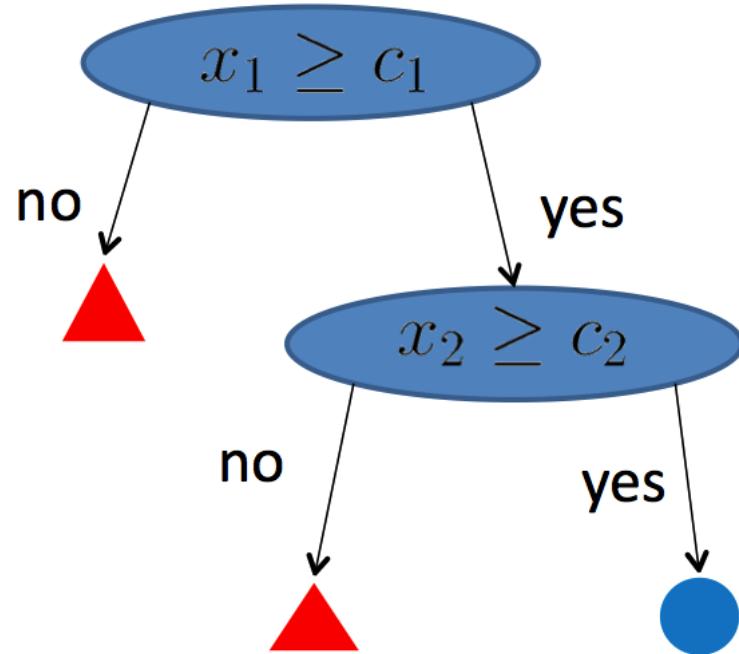


# DECISION TREES



“Stump” conditions:  $x > c$

# DECISION TREES



# DECISION TREE

- fast & intuitive prediction
- building optimal decision tree is NP complete
- building tree from root using greedy optimization

each time we split one leaf, finding optimal feature and threshold

- need criterion to select best splitting (feature, threshold)

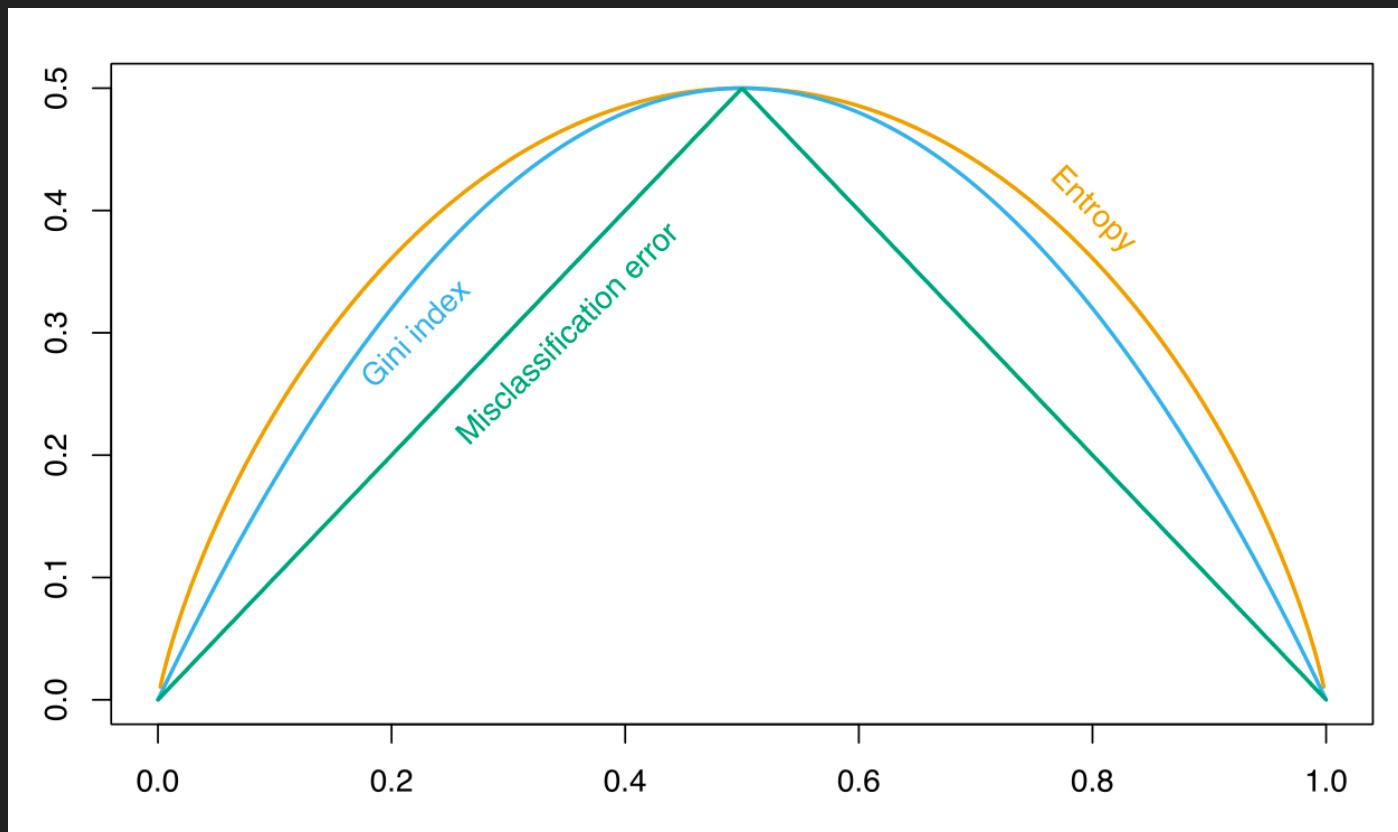
# SPLITTING CRITERIONS

$$\text{TotalImpurity} = \sum_{\text{leaf}} \text{impurity}(\text{leaf}) \times \text{size}(\text{leaf})$$

$$\text{Misclass.} = \min(p, 1 - p)$$

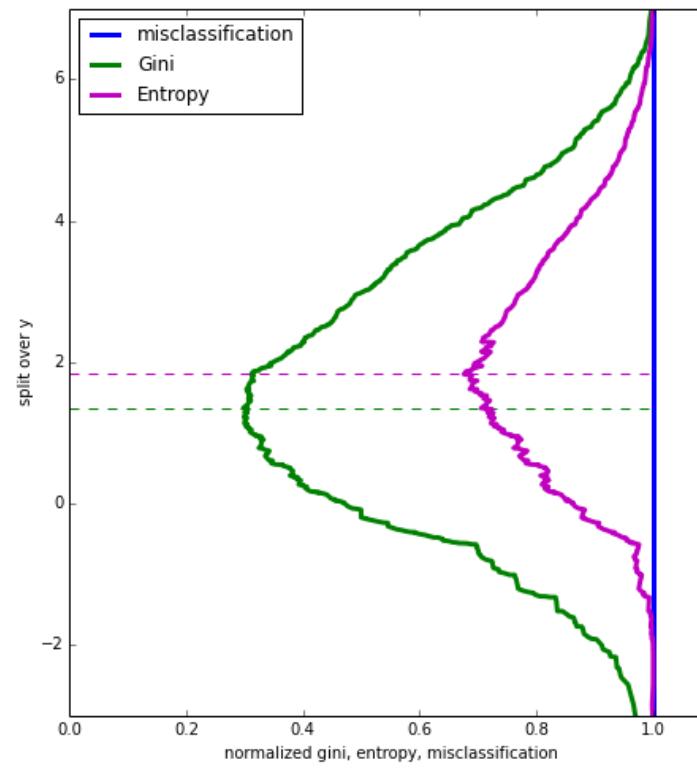
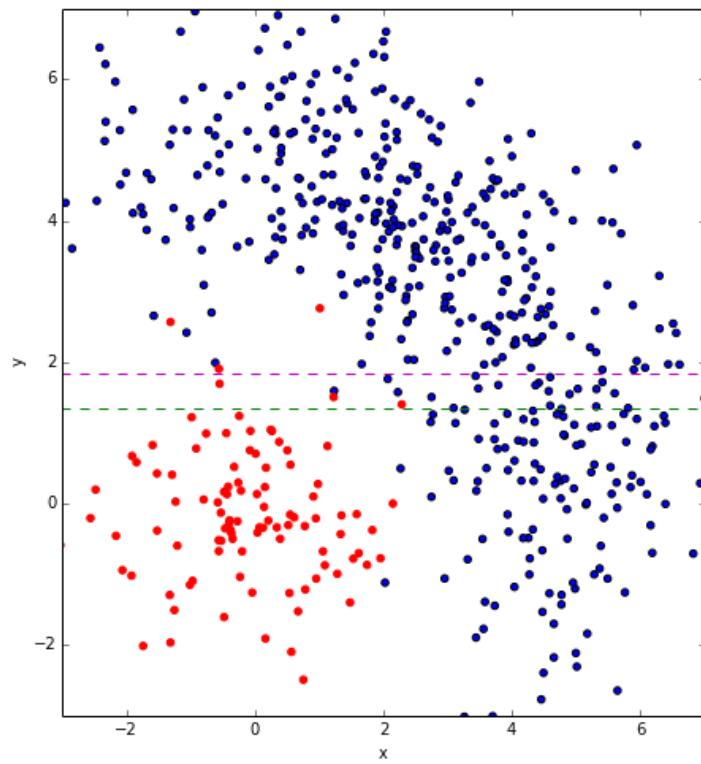
$$Gini = p(1 - p)$$

$$Entropy = -p \log p - (1 - p) \log(1 - p)$$



# SPLITTING CRITERIONS

Why using Gini or Entropy not misclassification?



# REGRESSION TREE

Greedy optimization (minimizing MSE):

$$\text{GlobalMSE} \sim \sum_i (y_i - \hat{y}_i)^2$$

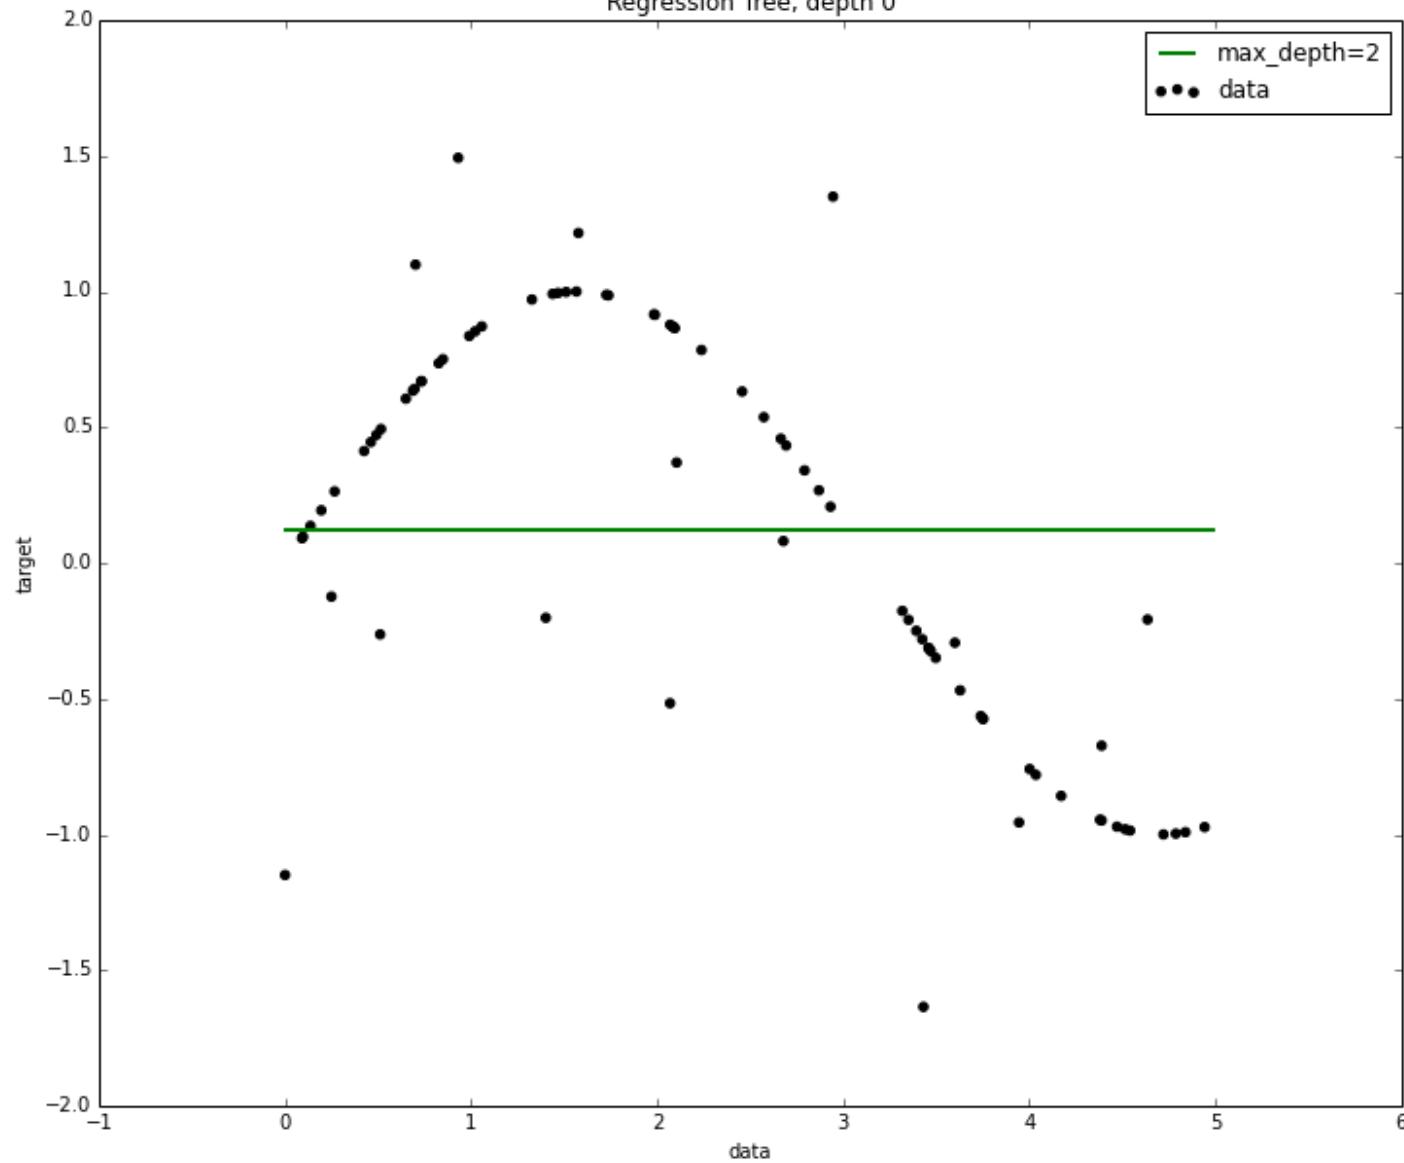
Can be rewritten as:

$$\text{GlobalMSE} \sim \sum_{\text{leaf}} \text{MSE}(\text{leaf}) \times \text{size}(\text{leaf})$$

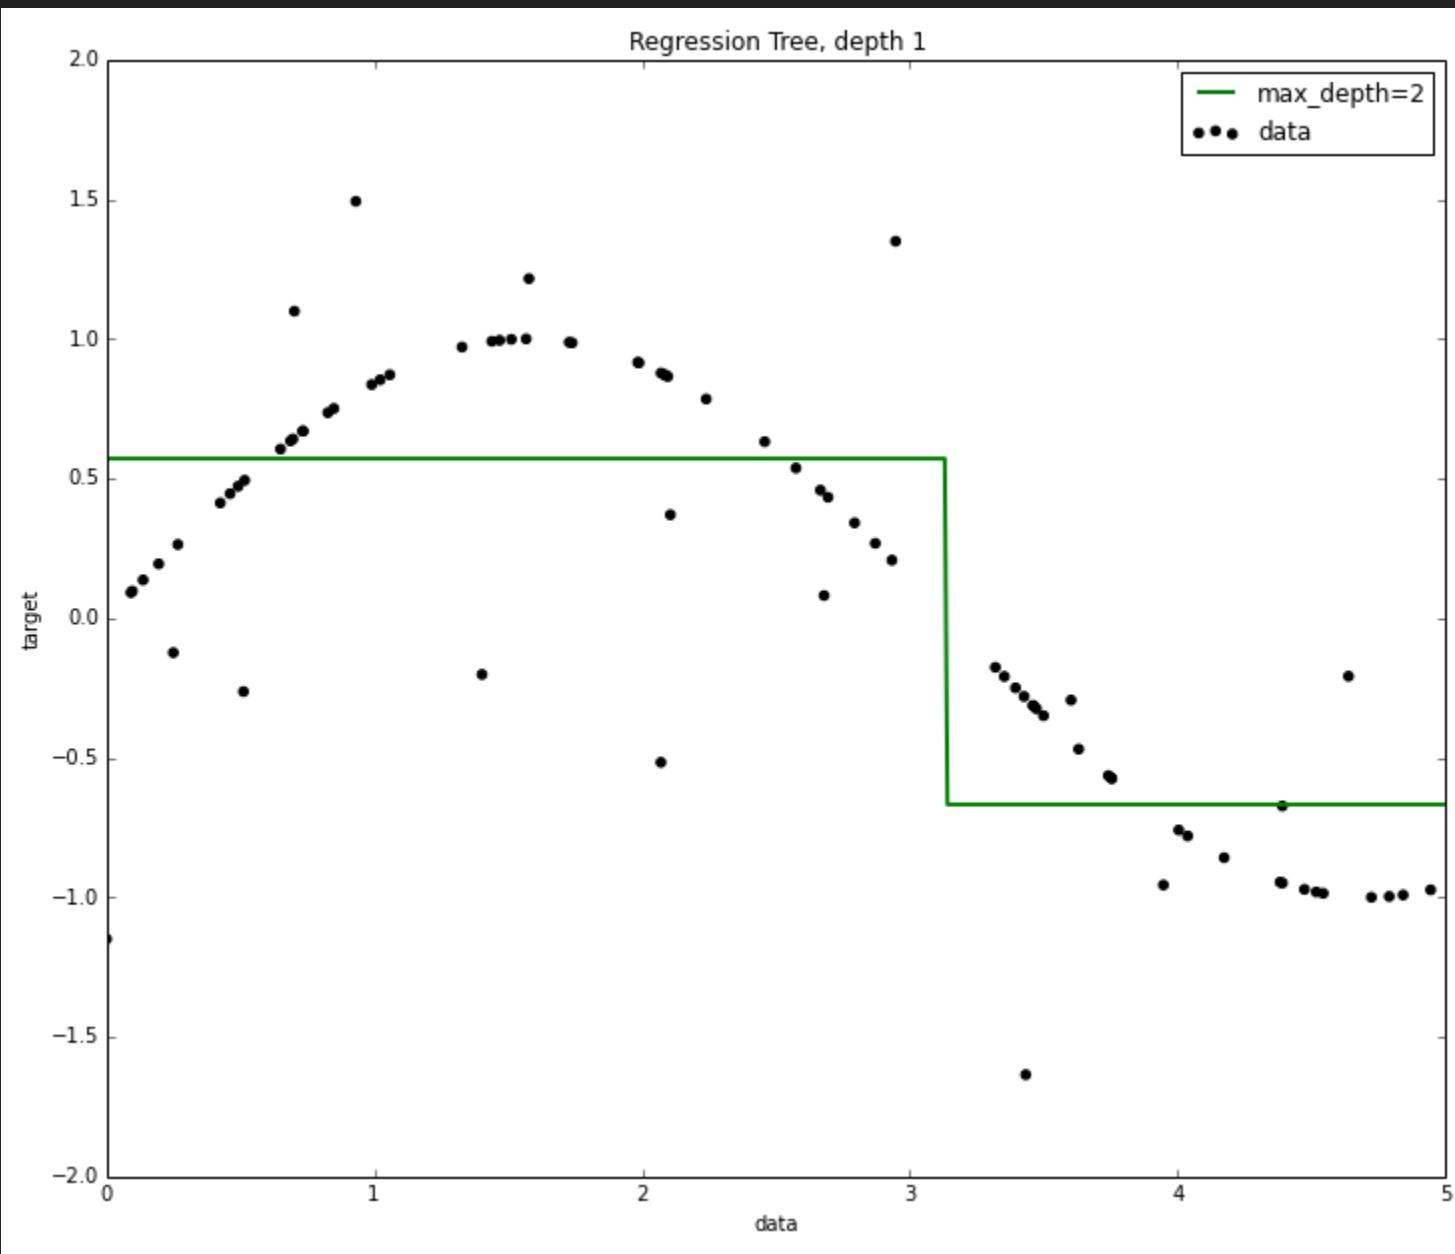
$\text{MSE}(\text{leaf})$  is like 'impurity' of leaf

$$\text{MSE}(\text{leaf}) = \frac{1}{\text{size}(\text{leaf})} \sum_{i \in \text{leaf}} (y_i - \hat{y}_i)^2$$

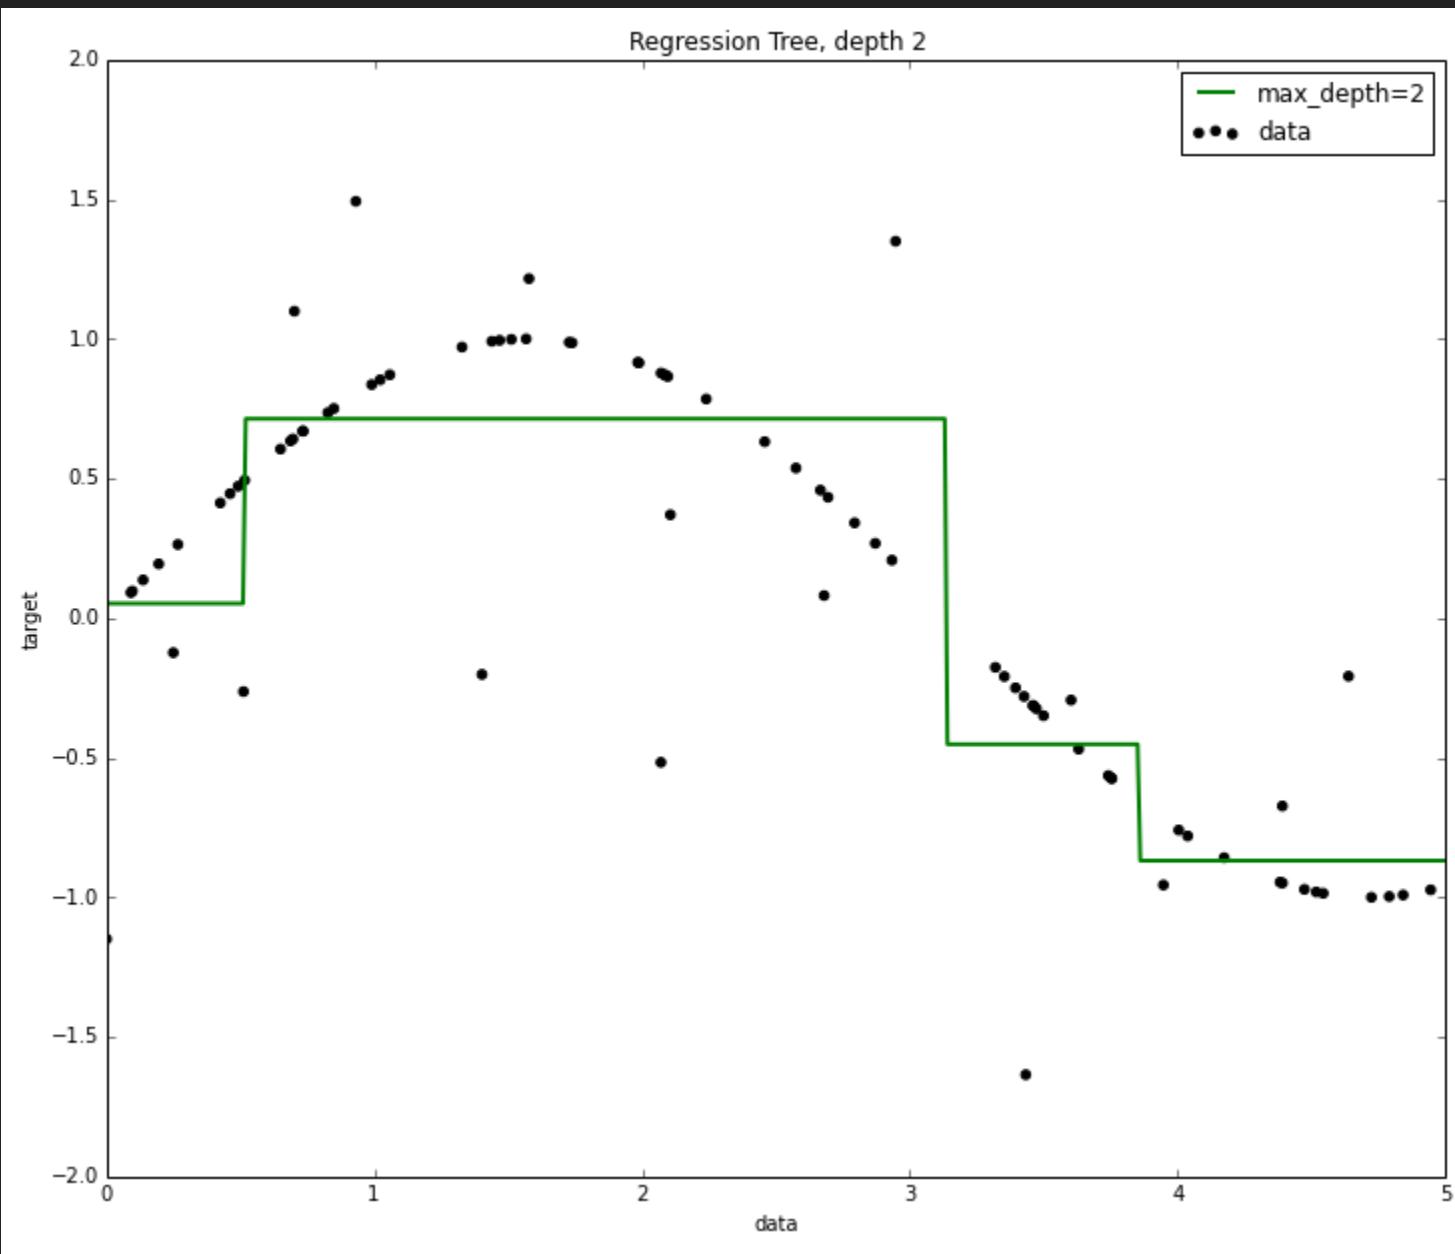
Regression Tree, depth 0



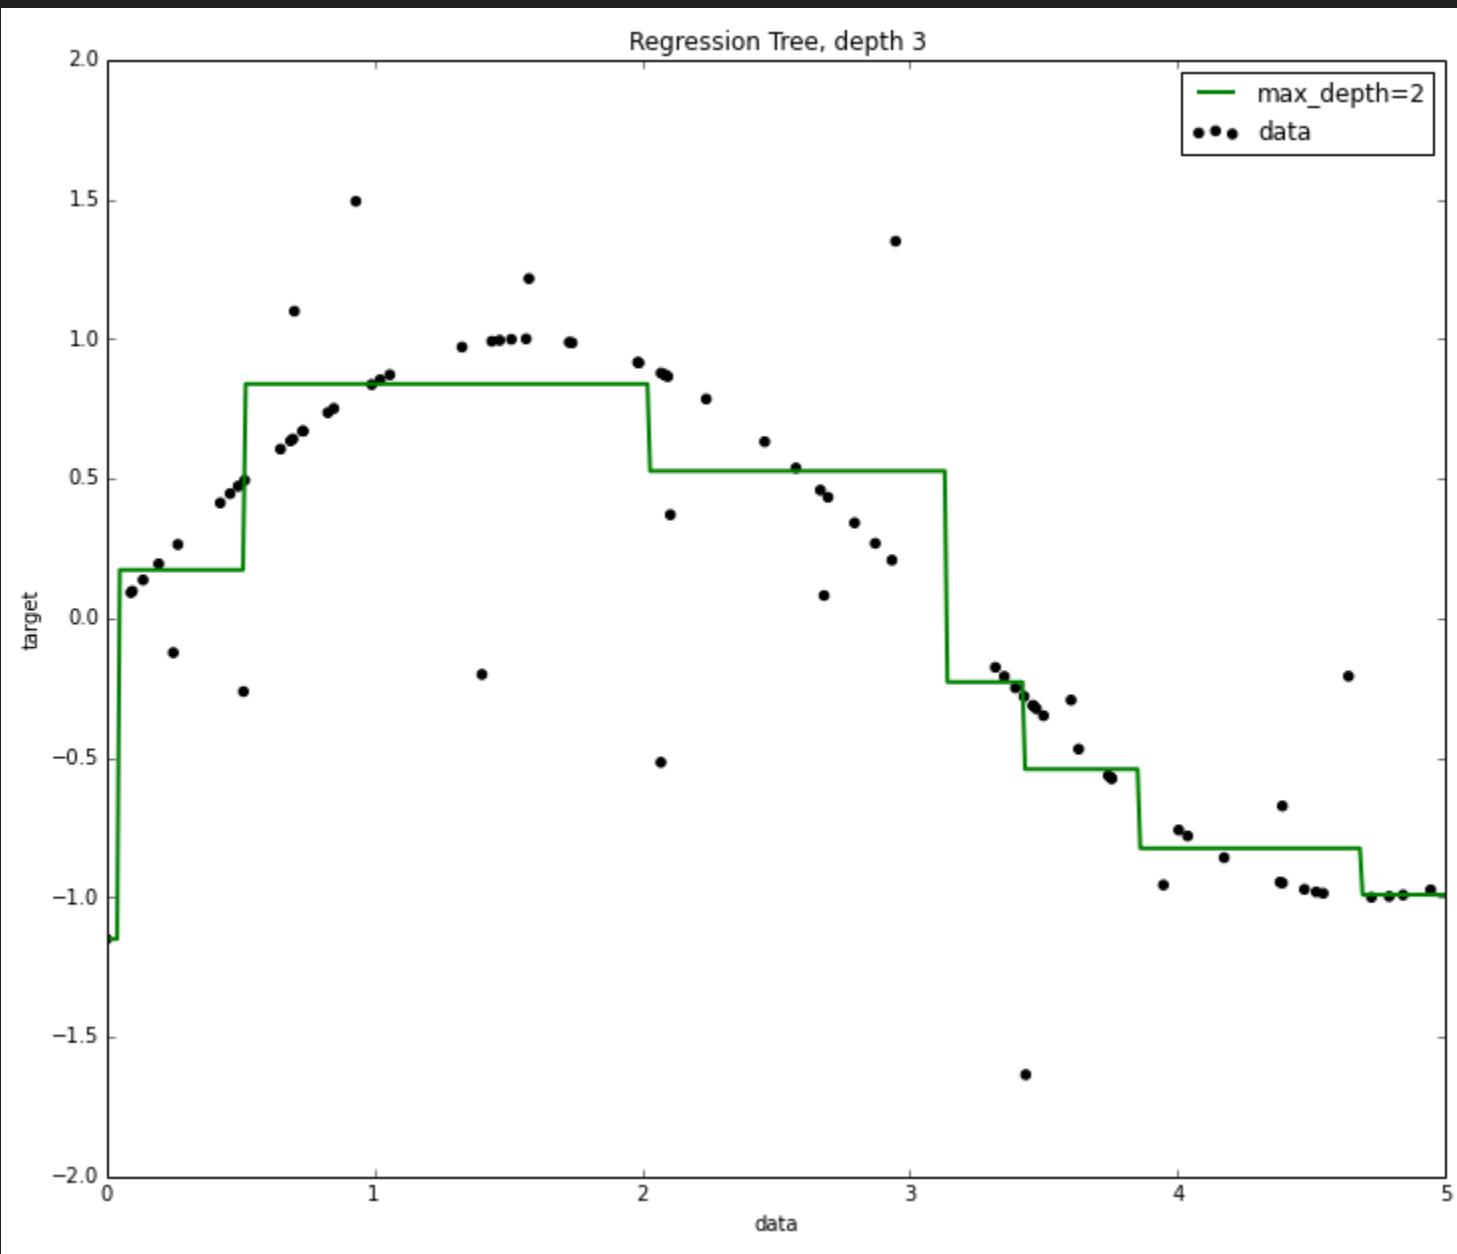
Regression Tree, depth 1



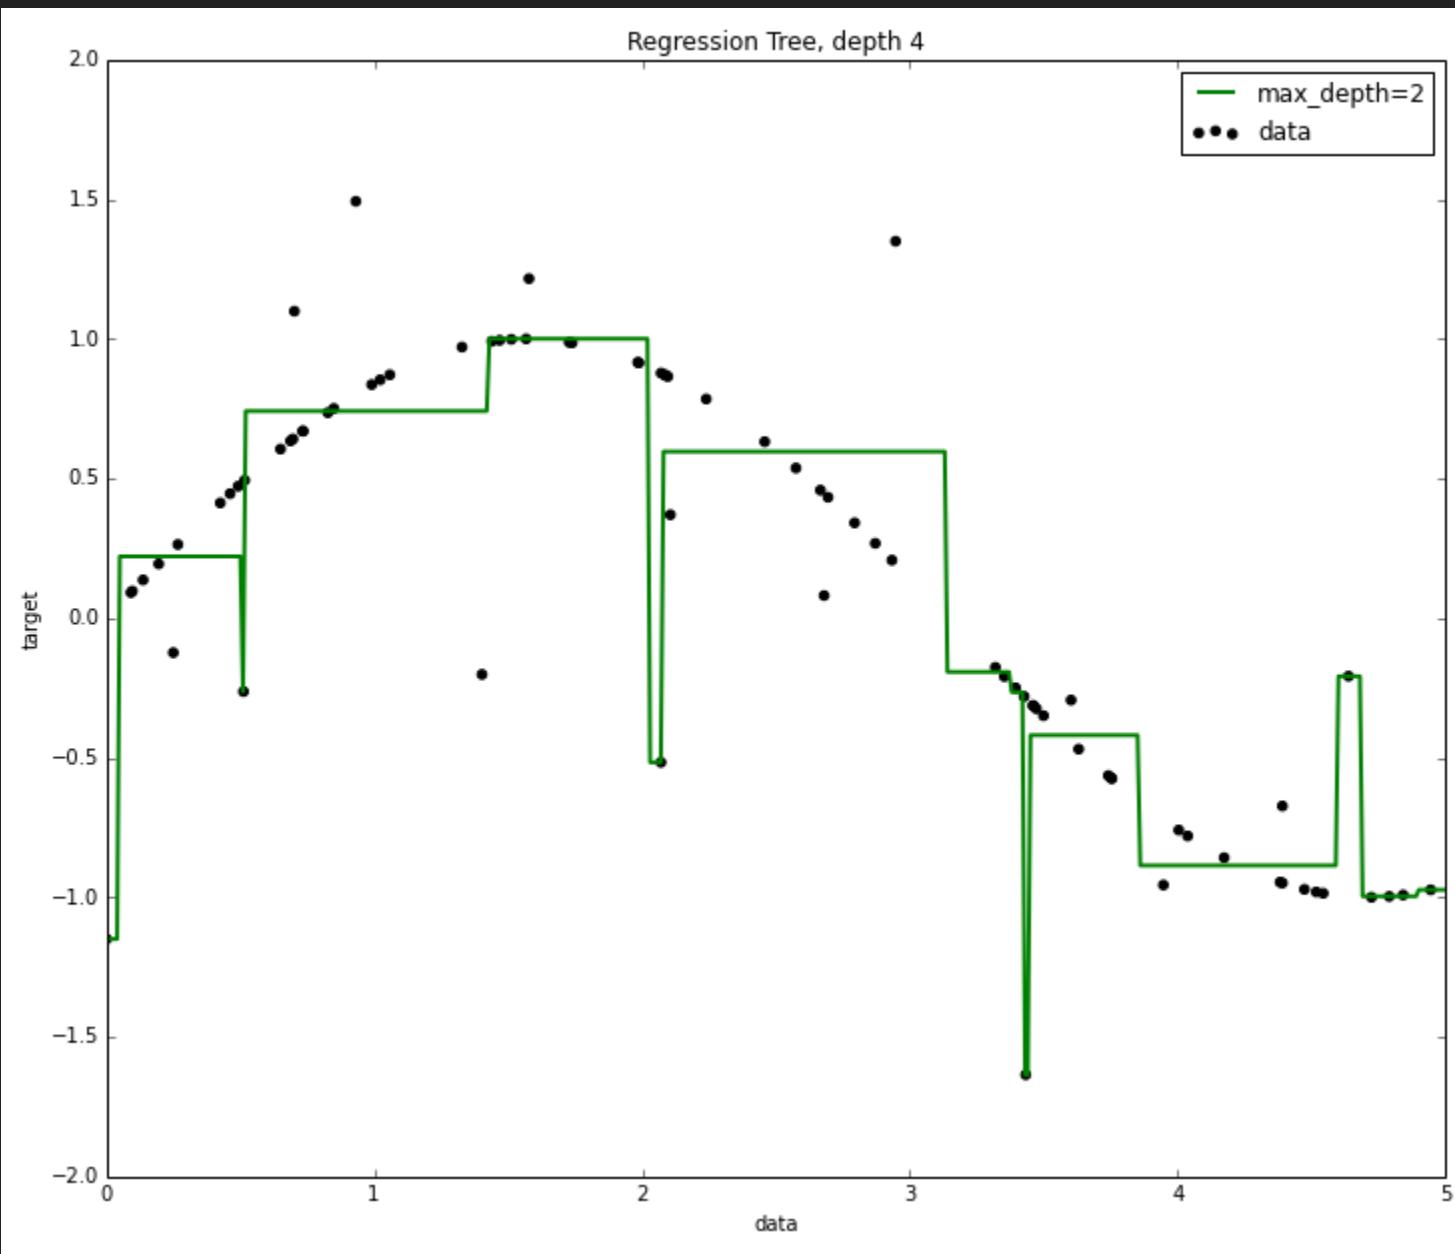
Regression Tree, depth 2



Regression Tree, depth 3



Regression Tree, depth 4



In most cases, regression trees are optimizing MSE:

$$\text{GlobalMSE} \sim \sum_i (y_i - \hat{y}_i)^2$$

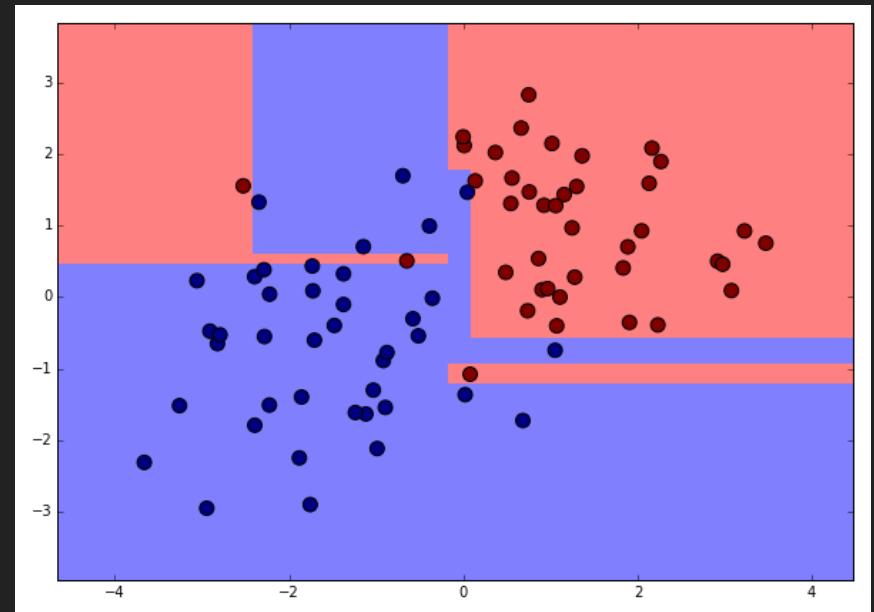
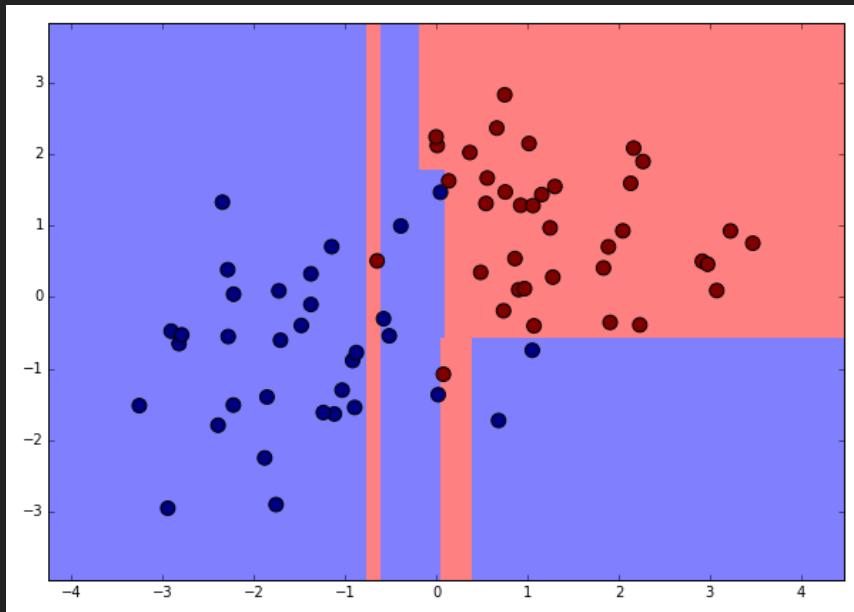
But other options also exist, i.e. MAE:

$$\text{GlobalMAE} \sim \sum_i |y_i - \hat{y}_i|$$

For MAE optimal value of leaf is median, not mean.

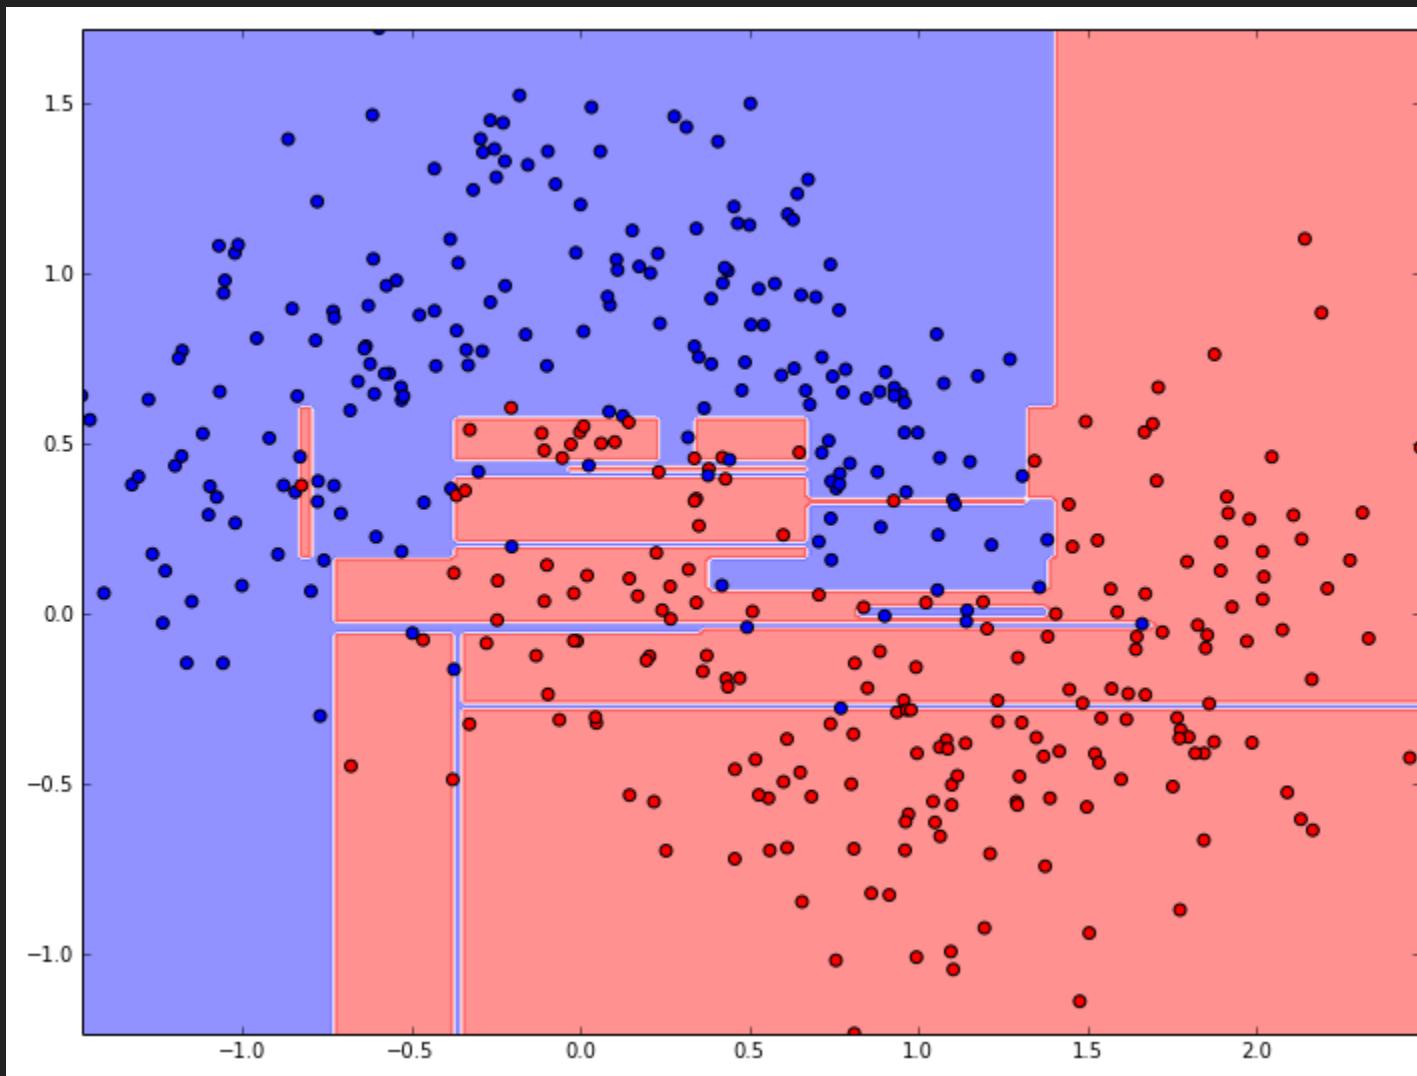
# DECISION TREES INSTABILITY

Little variation in training dataset produce different classification rule.



# PRE-STOPPING OF DECISION TREE

Tree keeps splitting until each event is correctly classified.

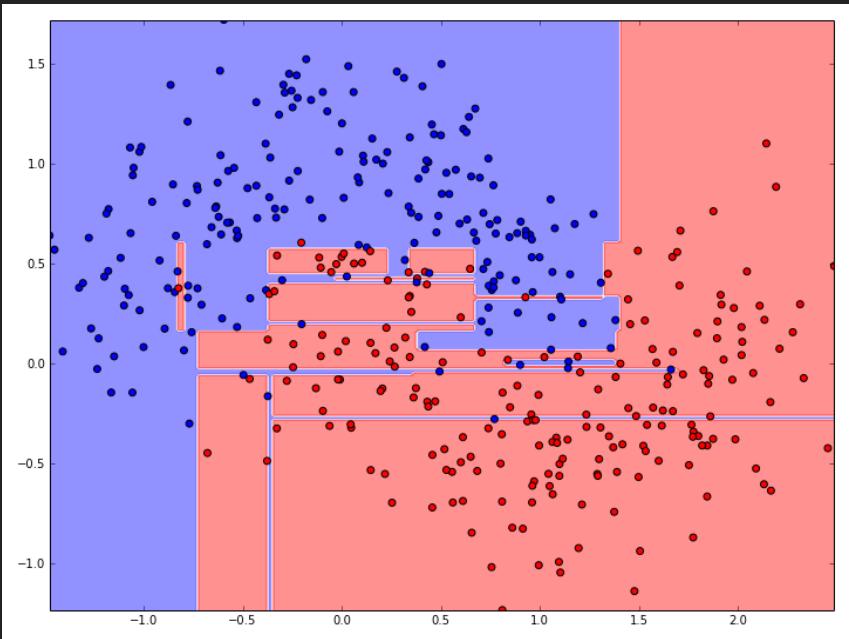


# PRE-STOPPING

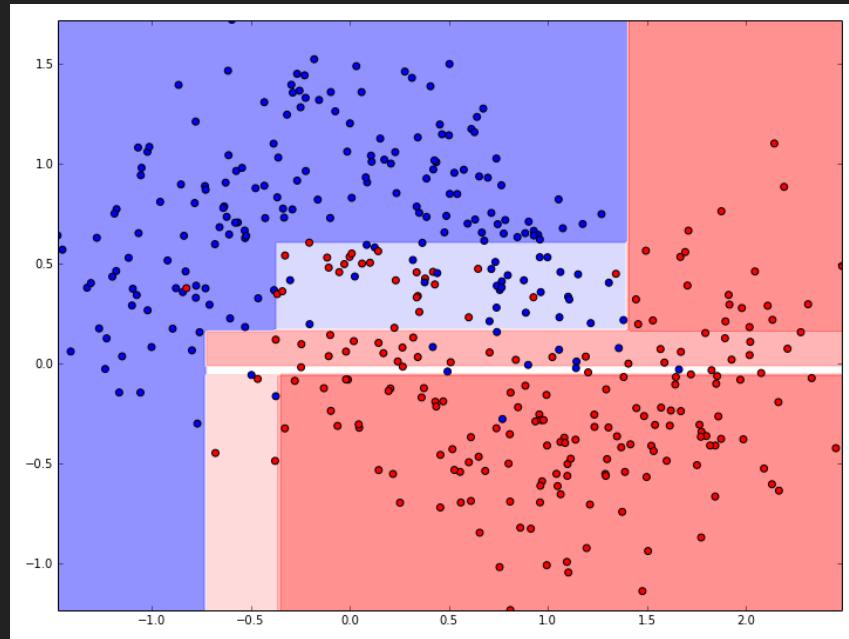
We can stop the process of splitting by imposing different restrictions.

- limit the depth of tree
- set minimal number of samples needed to split the leaf
- limit the minimal number of samples in leaf
- more advanced: maximal number of leaves in tree

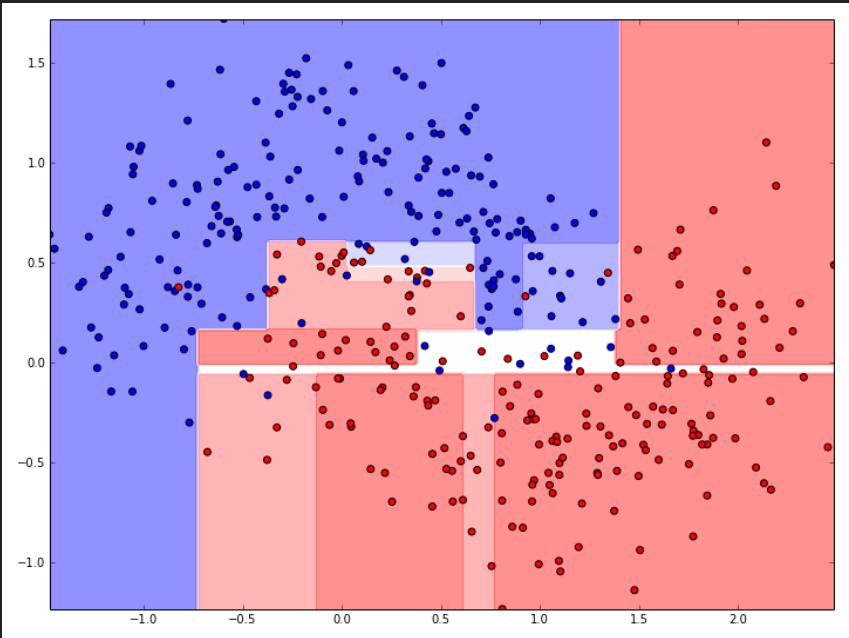
Any combinations of rules above is possible.



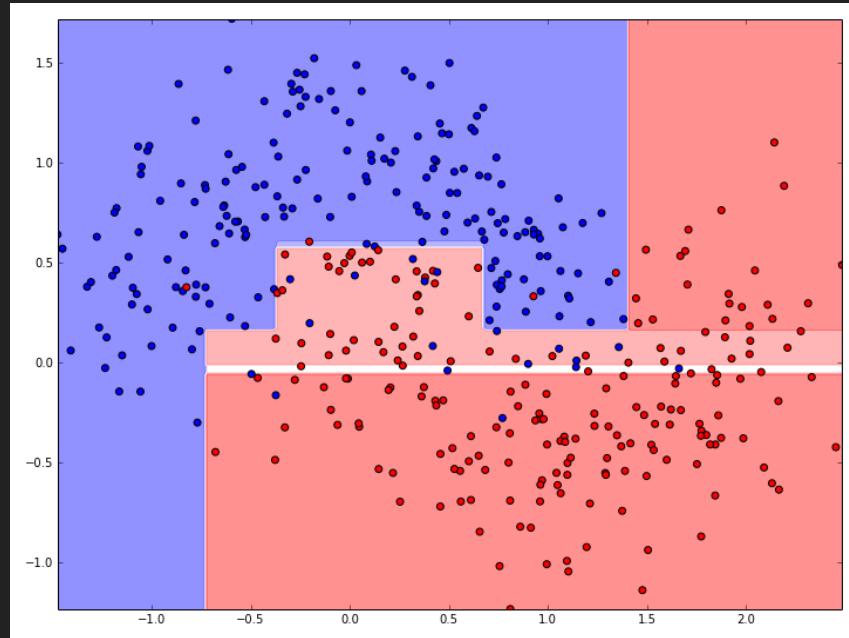
no prepruning



max\_depth



min # of samples in leaf



maximal number of leaves

# POST-PRUNING

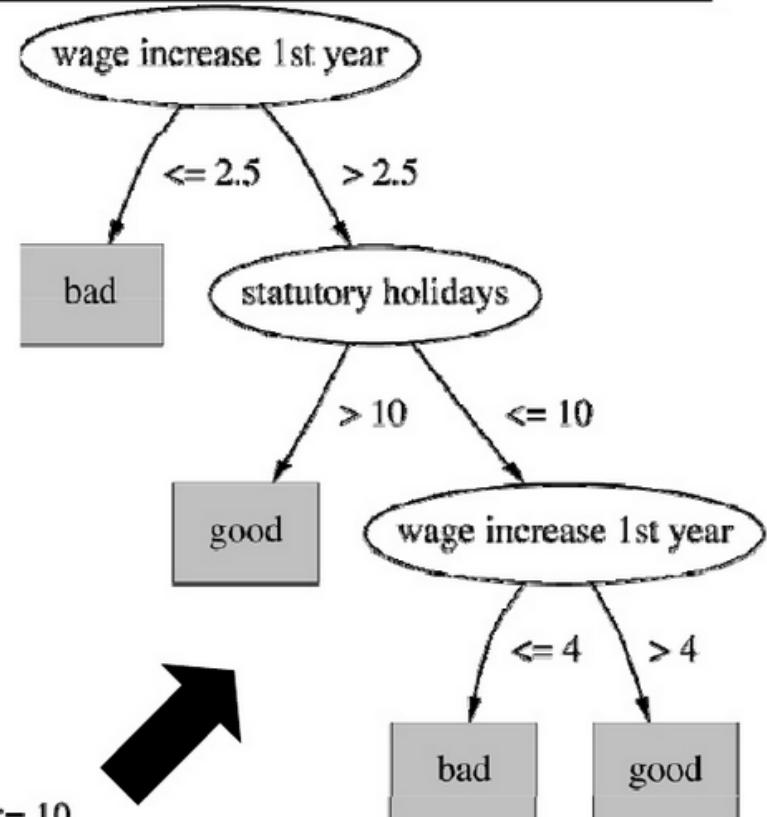
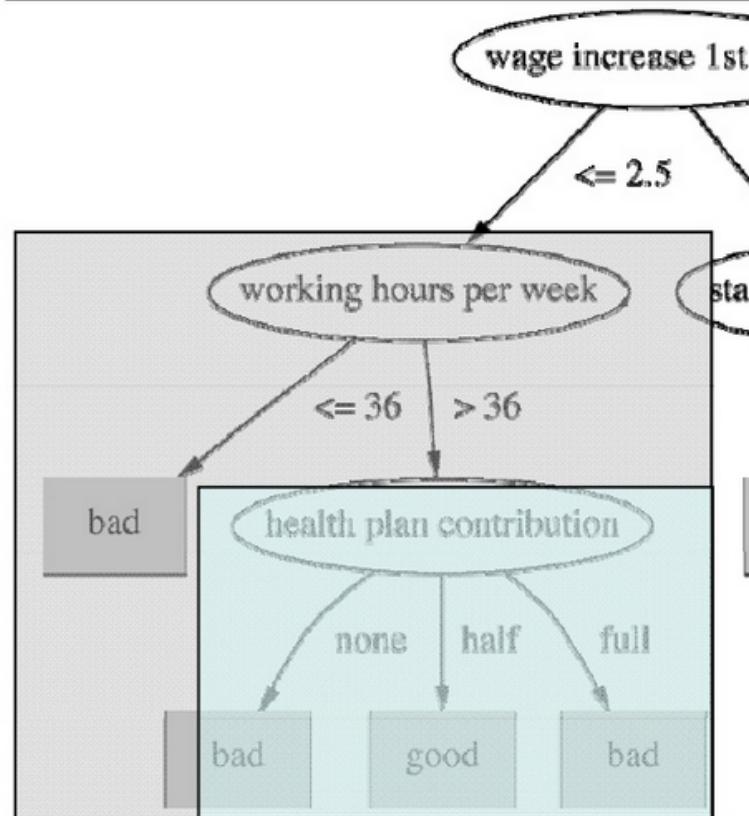
When tree tree is already built we can try optimize it to simplify formula.

Generally, much slower than pre-stopping.

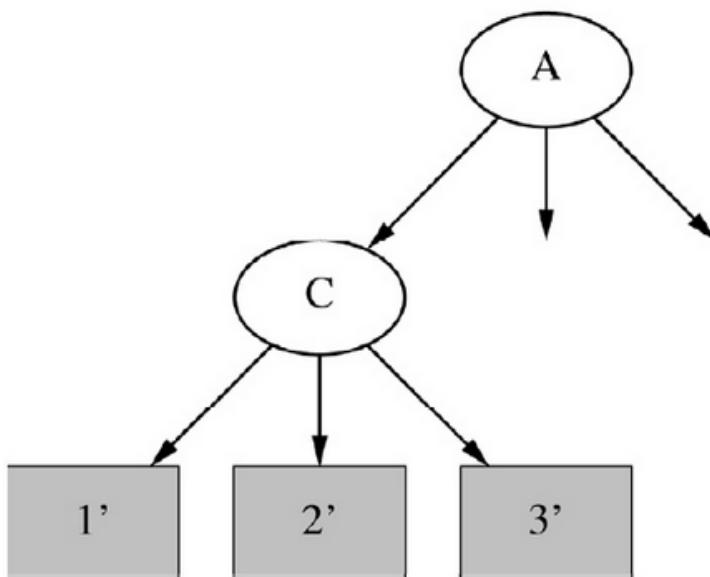
# Subtree Replacement

## Algorithm:

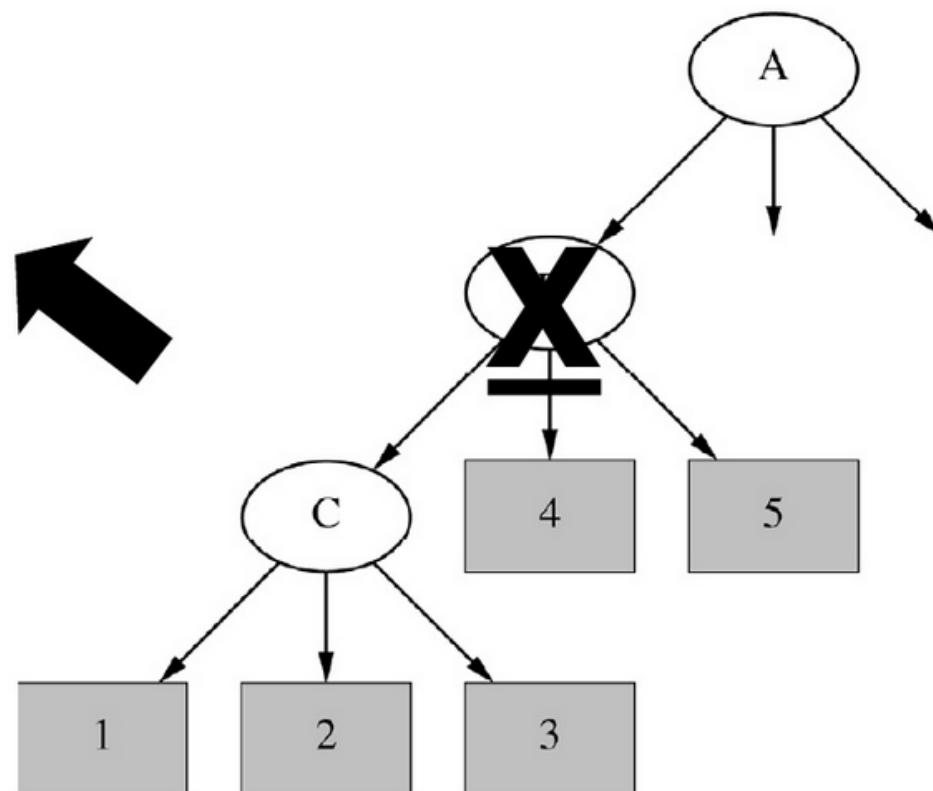
1. Split the data into training and validation set
2. Do until further pruning is harmful:
  - a. Evaluate impact on the validation set of pruning each possible node
  - b. Select the node whose removal most increases the validation set accuracy



# Subtree Raising



- Delete node
- Redistribute instances
- Slower than subtree replacement  
*(Worthwhile?)*



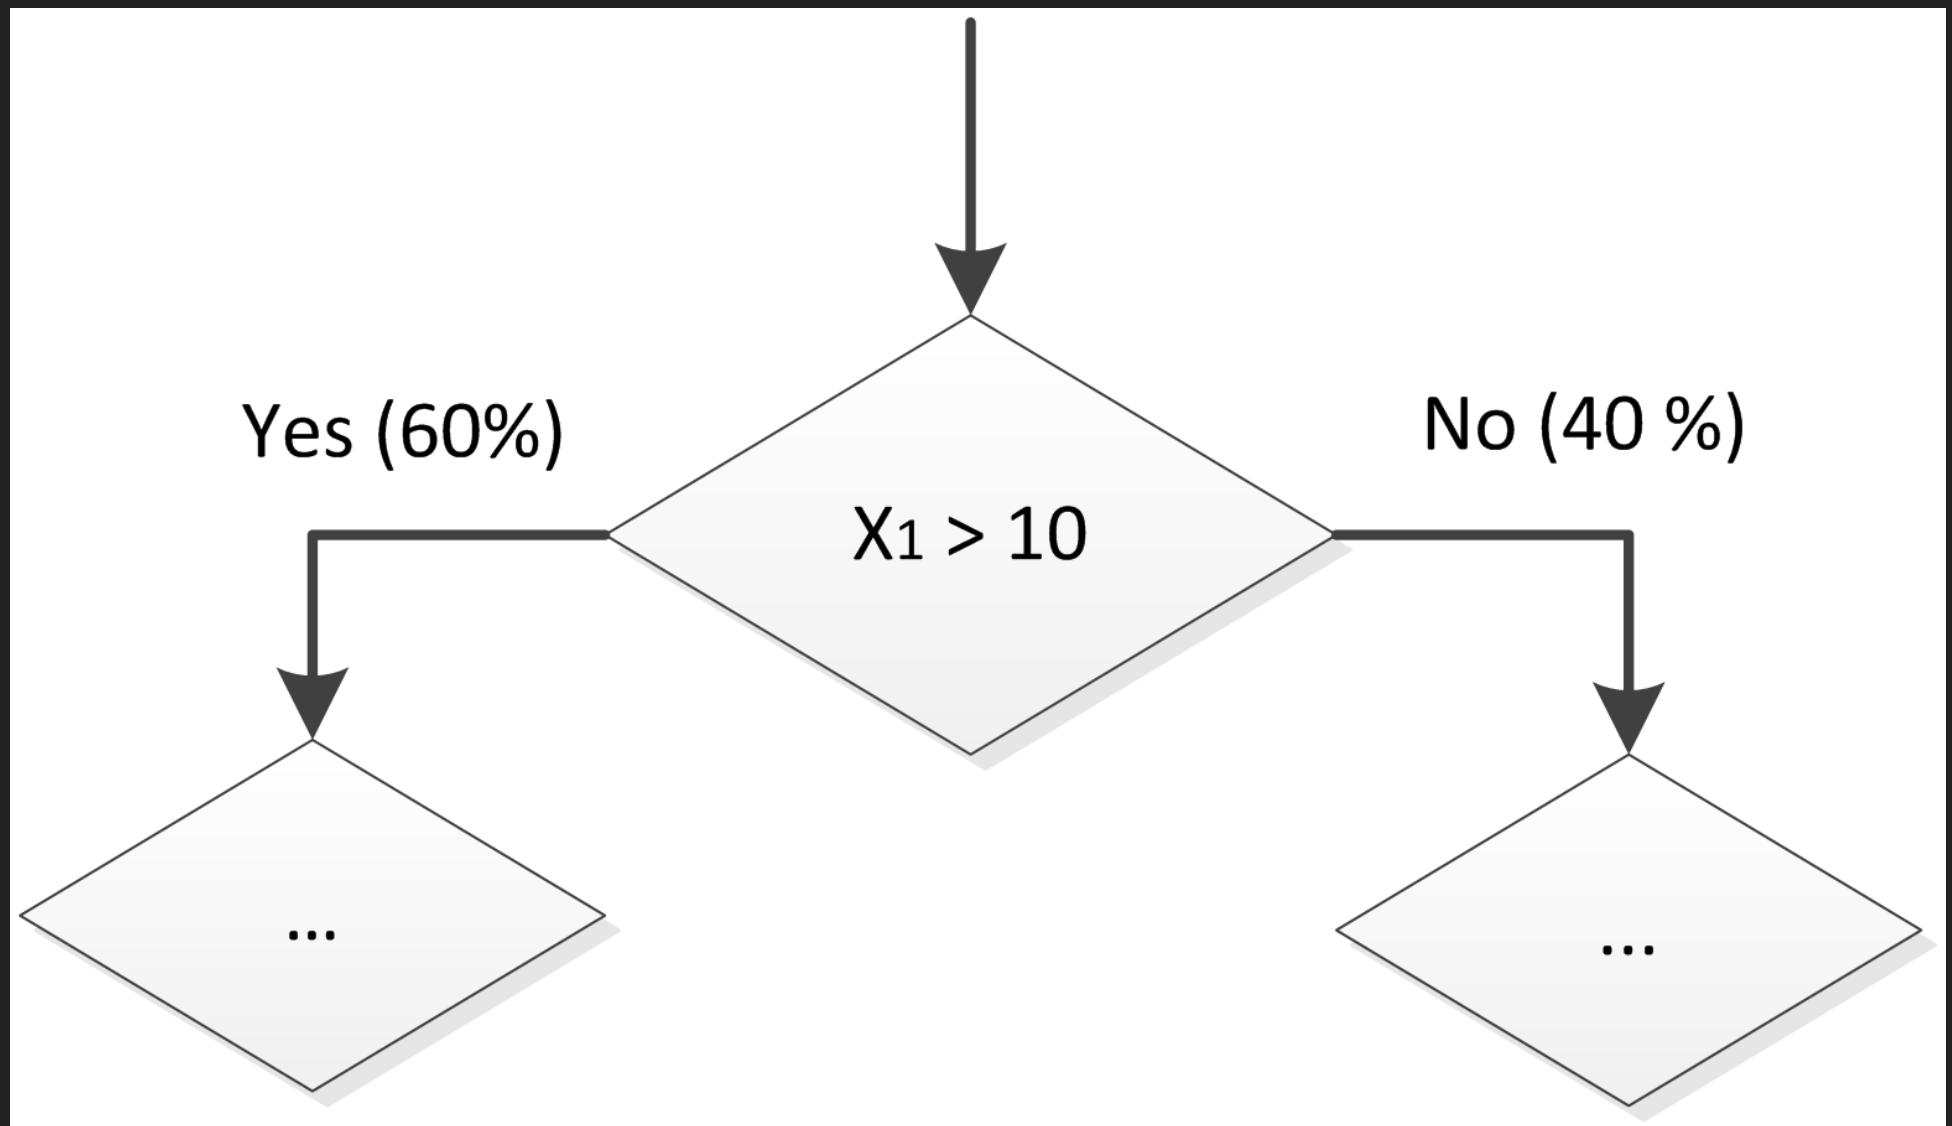
# SUMMARY OF DECISION TREE

1. Very intuitive algorithm for regression and classification
2. Fast prediction
3. Scale-independent
4. Supports multiclassification

But

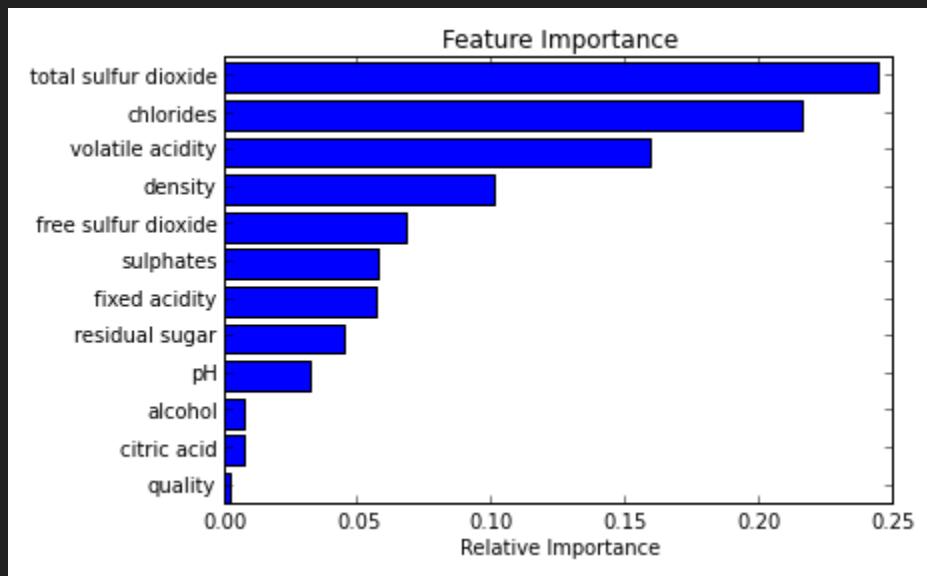
1. Training optimal tree is NP-complex
2. Trained greedily by optimizing Gini index or entropy (fast!)
3. Non-stable
4. Uses only trivial conditions

# MISSING VALUES IN DECISION TREES



If event being predicted lacks  $x_1$ , we use prior probabilities.

# FEATURE IMPORTANCES



Different approaches exist to measure importance of feature in final model

Importance of feature  $\neq$  quality provided by one feature

# FEATURE IMPORTANCES

- tree: counting number of splits made over this feature
- tree: counting gain in purity (e.g. Gini)  
fast and adequate
- common recipe: train without one feature,  
compare quality on test with/without one feature  
requires many evaluations
- common recipe: feature shuffling  
take one column in test dataset and shuffle them. Compare  
quality with/without shuffling.

# THE END

Tomorrow: ensembles and boosting